

Databases

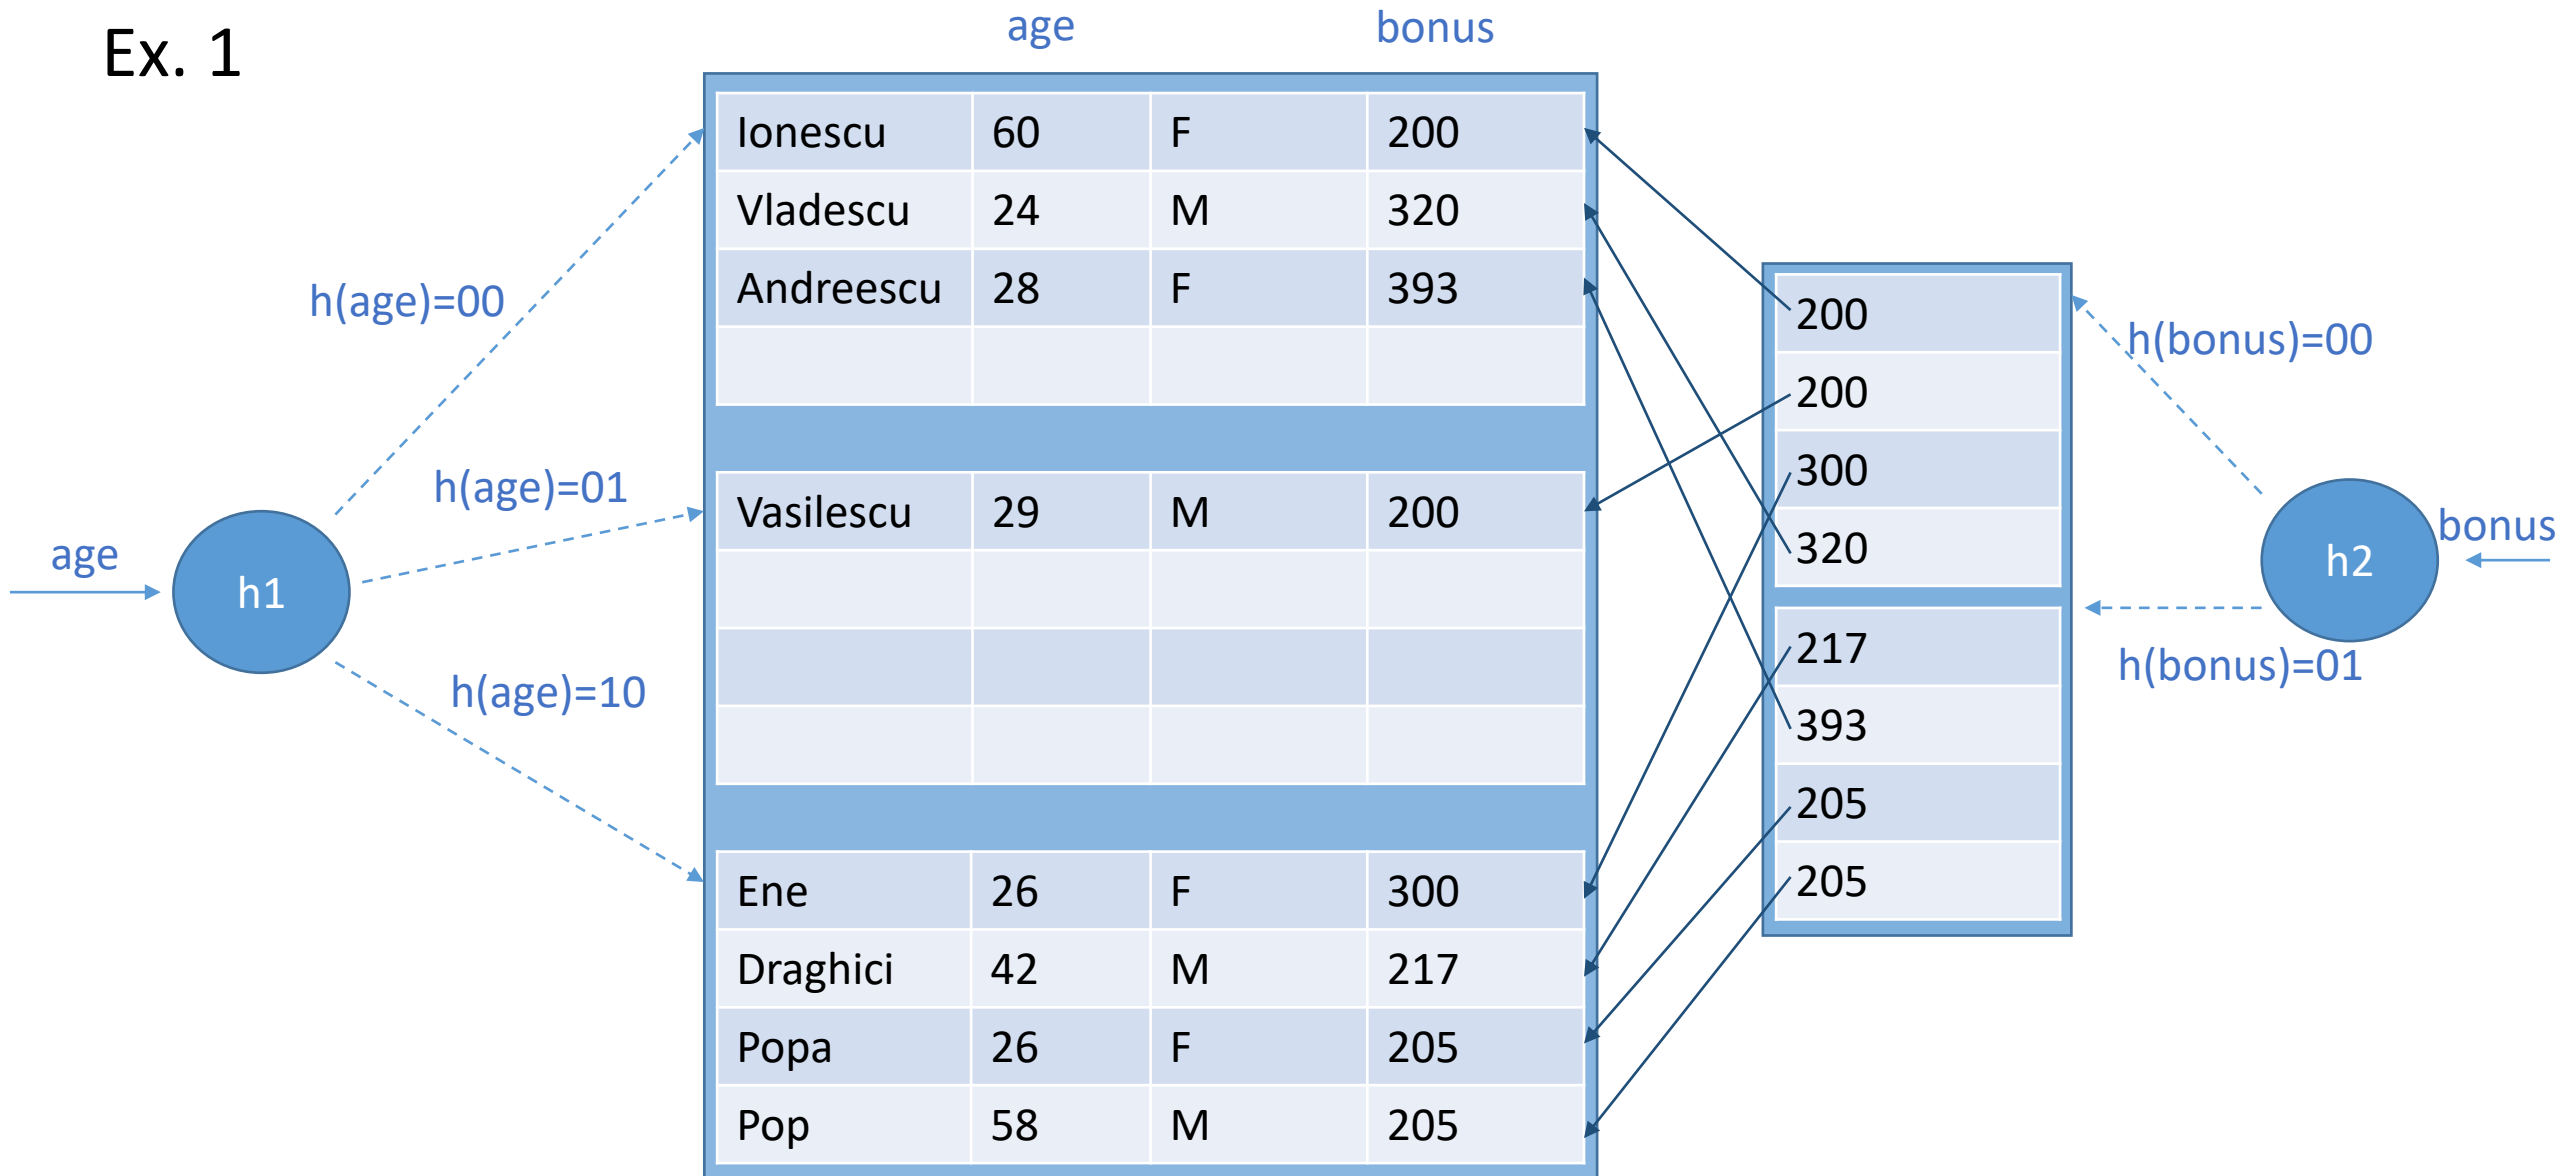
Lecture 9

The Physical Structure of Databases (II)

- Indexes. Tree-Structured Indexing -

Indexes - Data Entries

Ex. 1



Indexes - Data Entries

Ex. 1

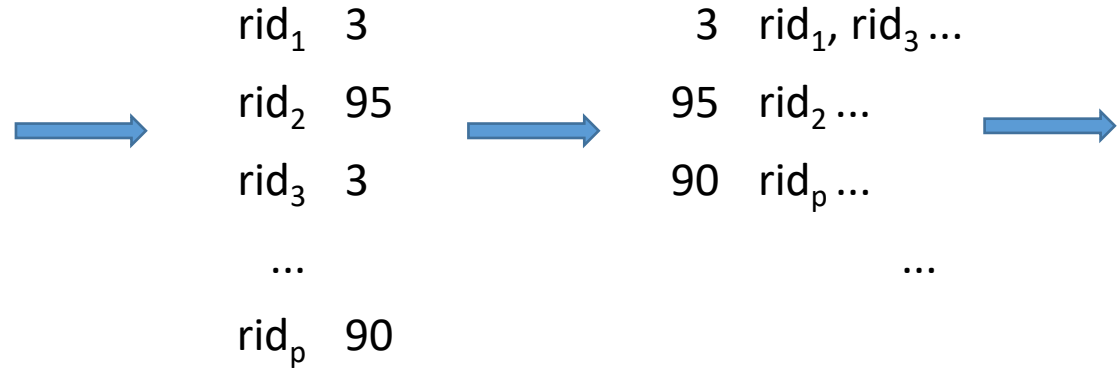
- file with Employee records hashed on *age*
 - record <Ionescu, 60, F, 200>:
 - apply hash function to *age*: convert 60 to its binary representation, take the 2 least significant bits as the bucket identifier for the record
- index file that uses alternative 1 (data entries are the actual data records), search key *age*
- index that uses alternative 2 (data entries have the form <search key, rid>), search key *bonus*
- both indexes use hashing to locate data entries

Indexes - Data Entries

Ex. 2

	Name	Score	Age
rid ₁	Popescu	3	44
rid ₂	Ionescu	95	80
rid ₃	Vladescu	3	45
...		...	
rid _p	Xulescu	90	14

search key



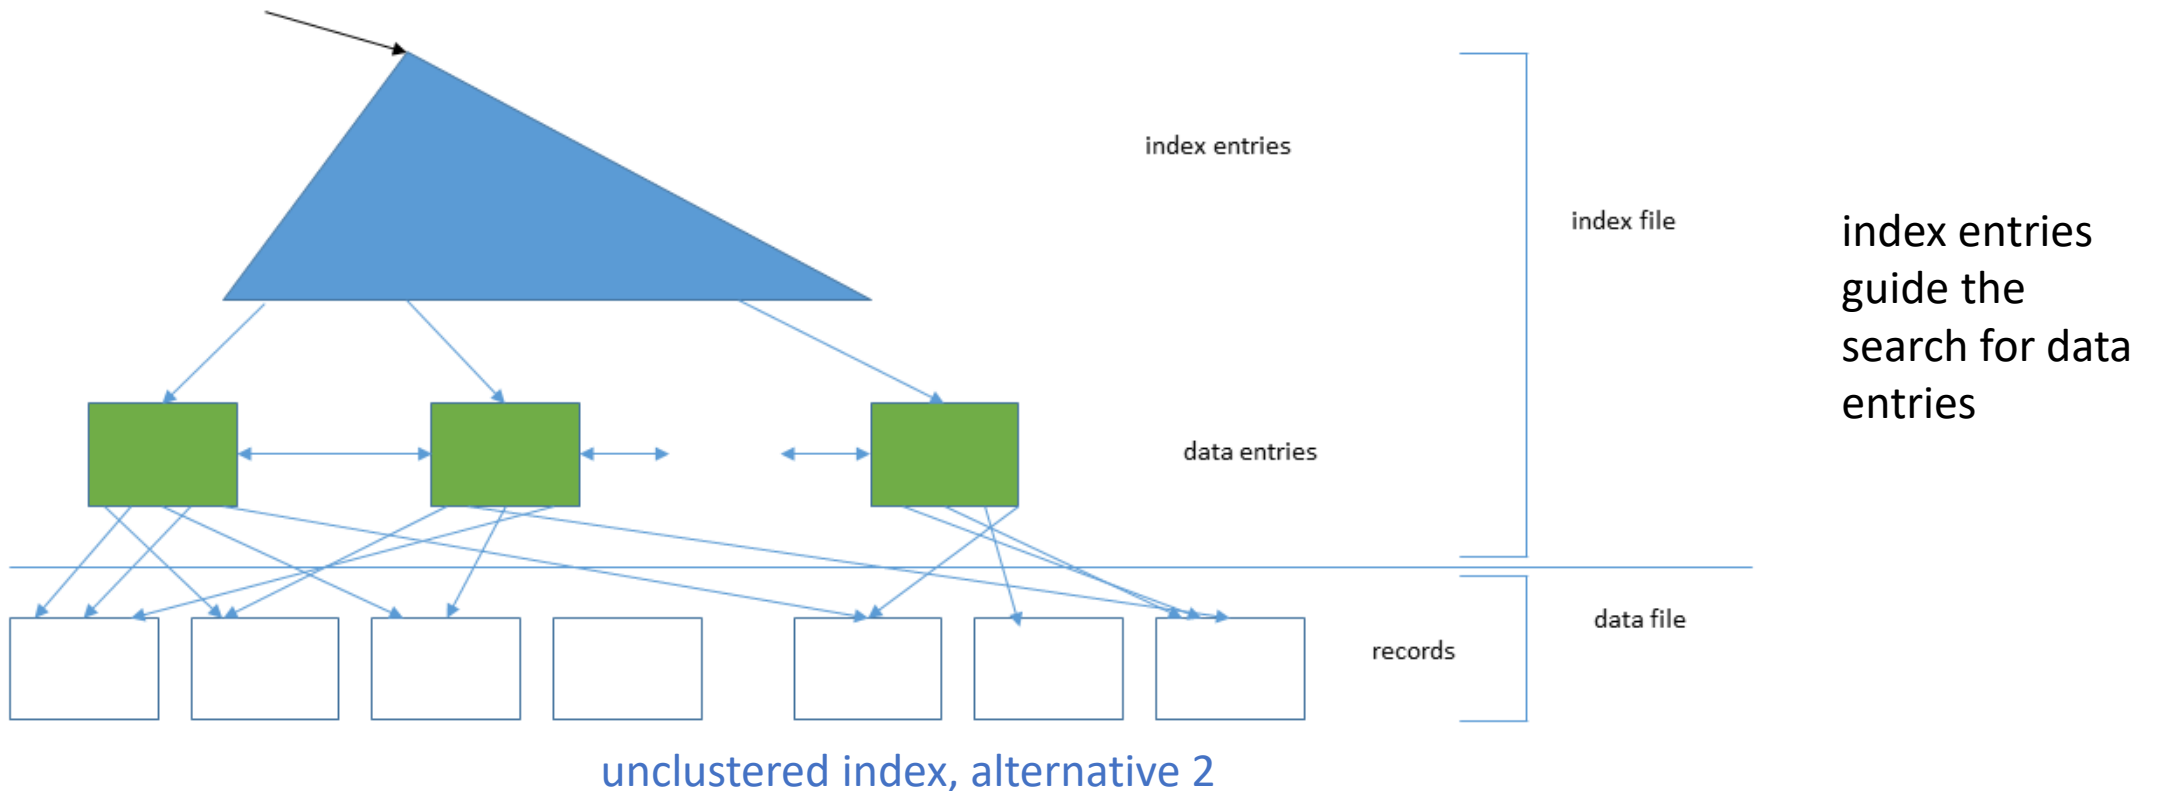
An arrow labeled 'data entry' points to the first row of an index file table. The table has two columns: the first contains scores (3, 90, 95, ...) and the second contains the corresponding rid values (rid₁, rid₃..., rid_p..., rid₂..., ...).

...	
3	rid ₁ , rid ₃ ...
90	rid _p ...
95	rid ₂ ...
...	

index file

Clustered / Unclustered Indexes

- clustered index: the order of the data records is close to / the same as the order of the data entries
- unclustered index: index that is not clustered



Clustered / Unclustered Indexes

- index that uses alternative 1 - clustered (by definition, since the data entries are the actual data records)
- indexes using alternatives 2 / 3 are clustered only if the data records are ordered on the search key
- in practice:
 - expensive to maintain the sort order for files, so they are rarely kept sorted
 - a clustered index is an index that uses alternative 1 for data entries
 - an index that uses alternative 2 or 3 for data entries is unclustered
- on a collection of records:
 - there can be at most 1 clustered index
 - and several unclustered indexes

Clustered / Unclustered Indexes

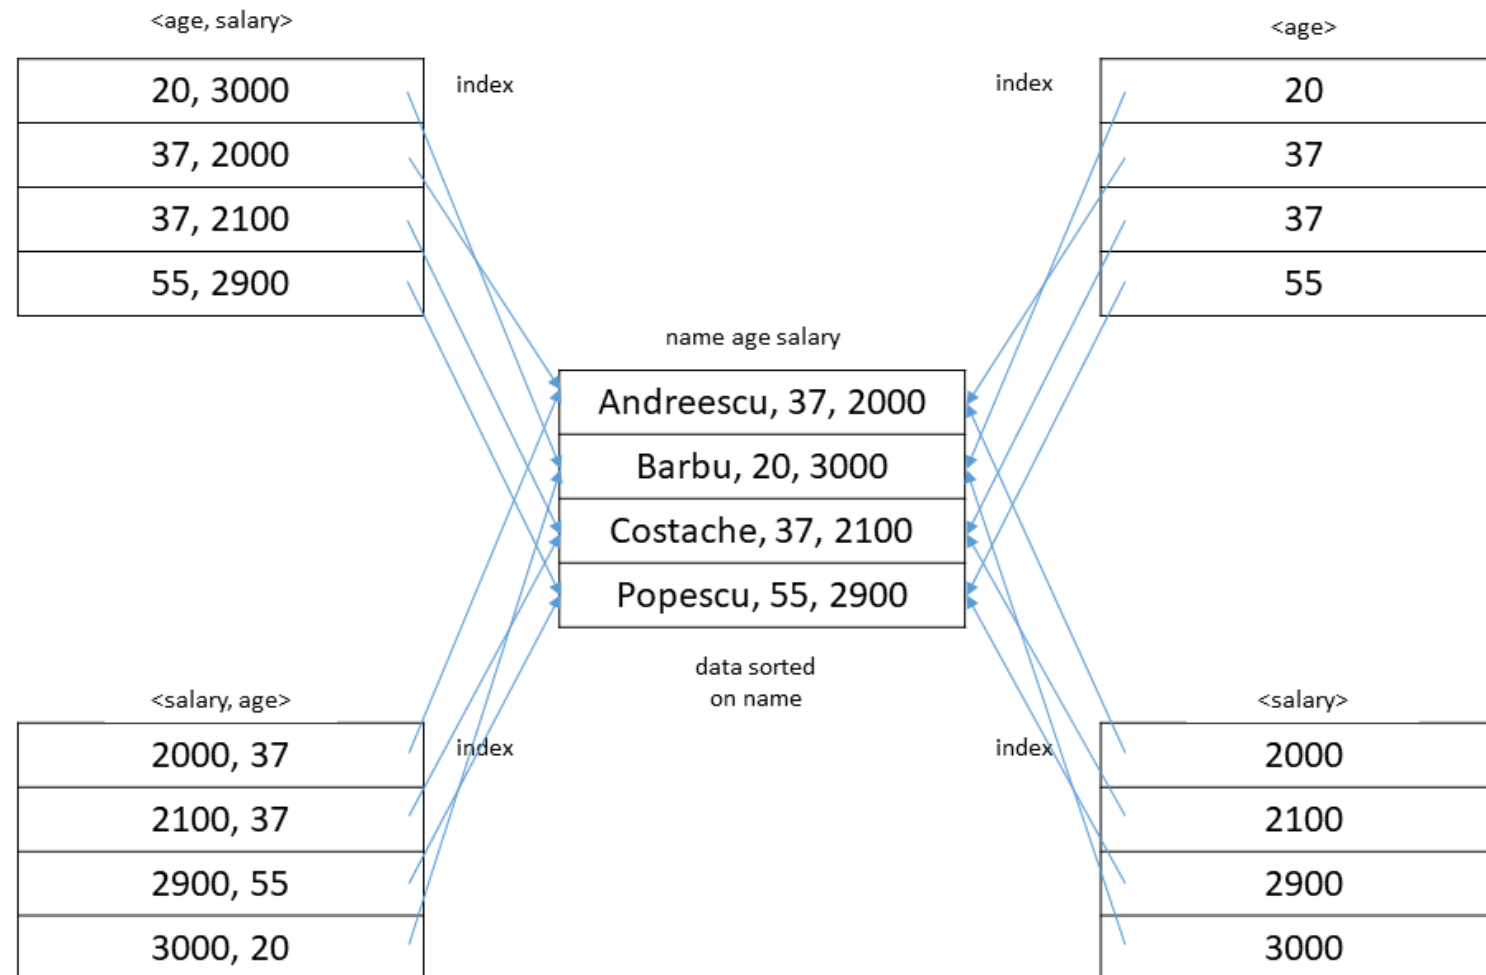
- range search query (e.g., *where age between 20 and 30*)
 - cost of using an unclustered index
 - each data entry that meets the condition in the query could contain a rid pointing to a distinct page
 - the number of I/O operations could be equal to the number of data entries that satisfy the query's condition

Primary / Secondary Indexes

- primary index
 - the search key includes the primary key
- secondary index
 - index that is not primary
- unique index
 - the search key contains a candidate key
- duplicates
 - data entries with the same search key value
- primary indexes, unique indexes cannot contain duplicates
- secondary indexes can contain duplicates

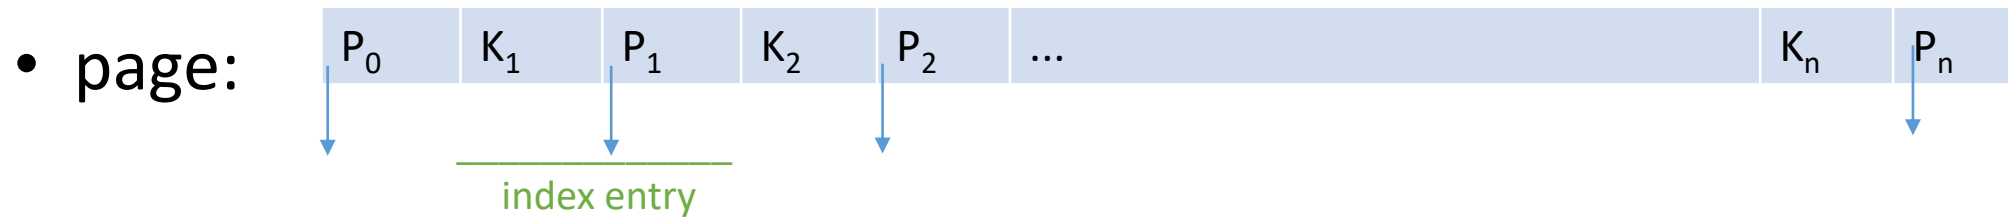
Composite Search Keys

- composite (concatenated) search key - search key that contains several fields
 - examples

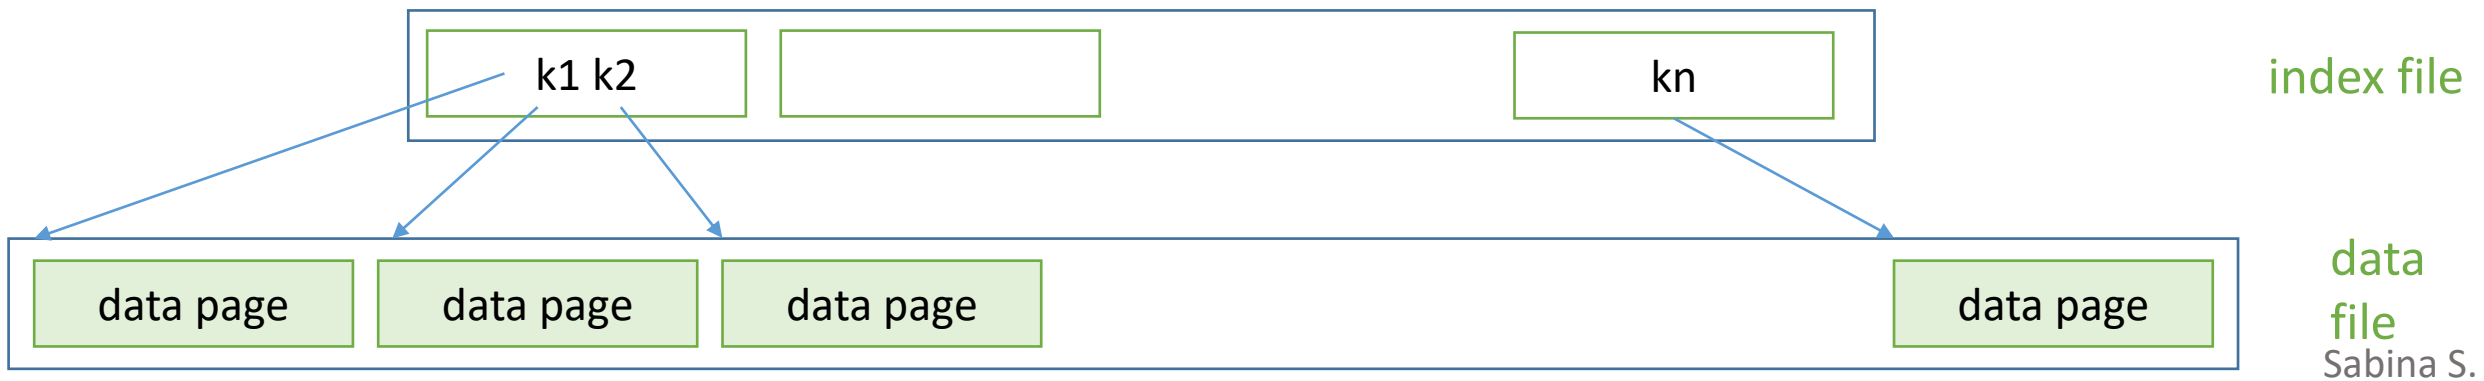


Indexed Sequential Access Method (ISAM)

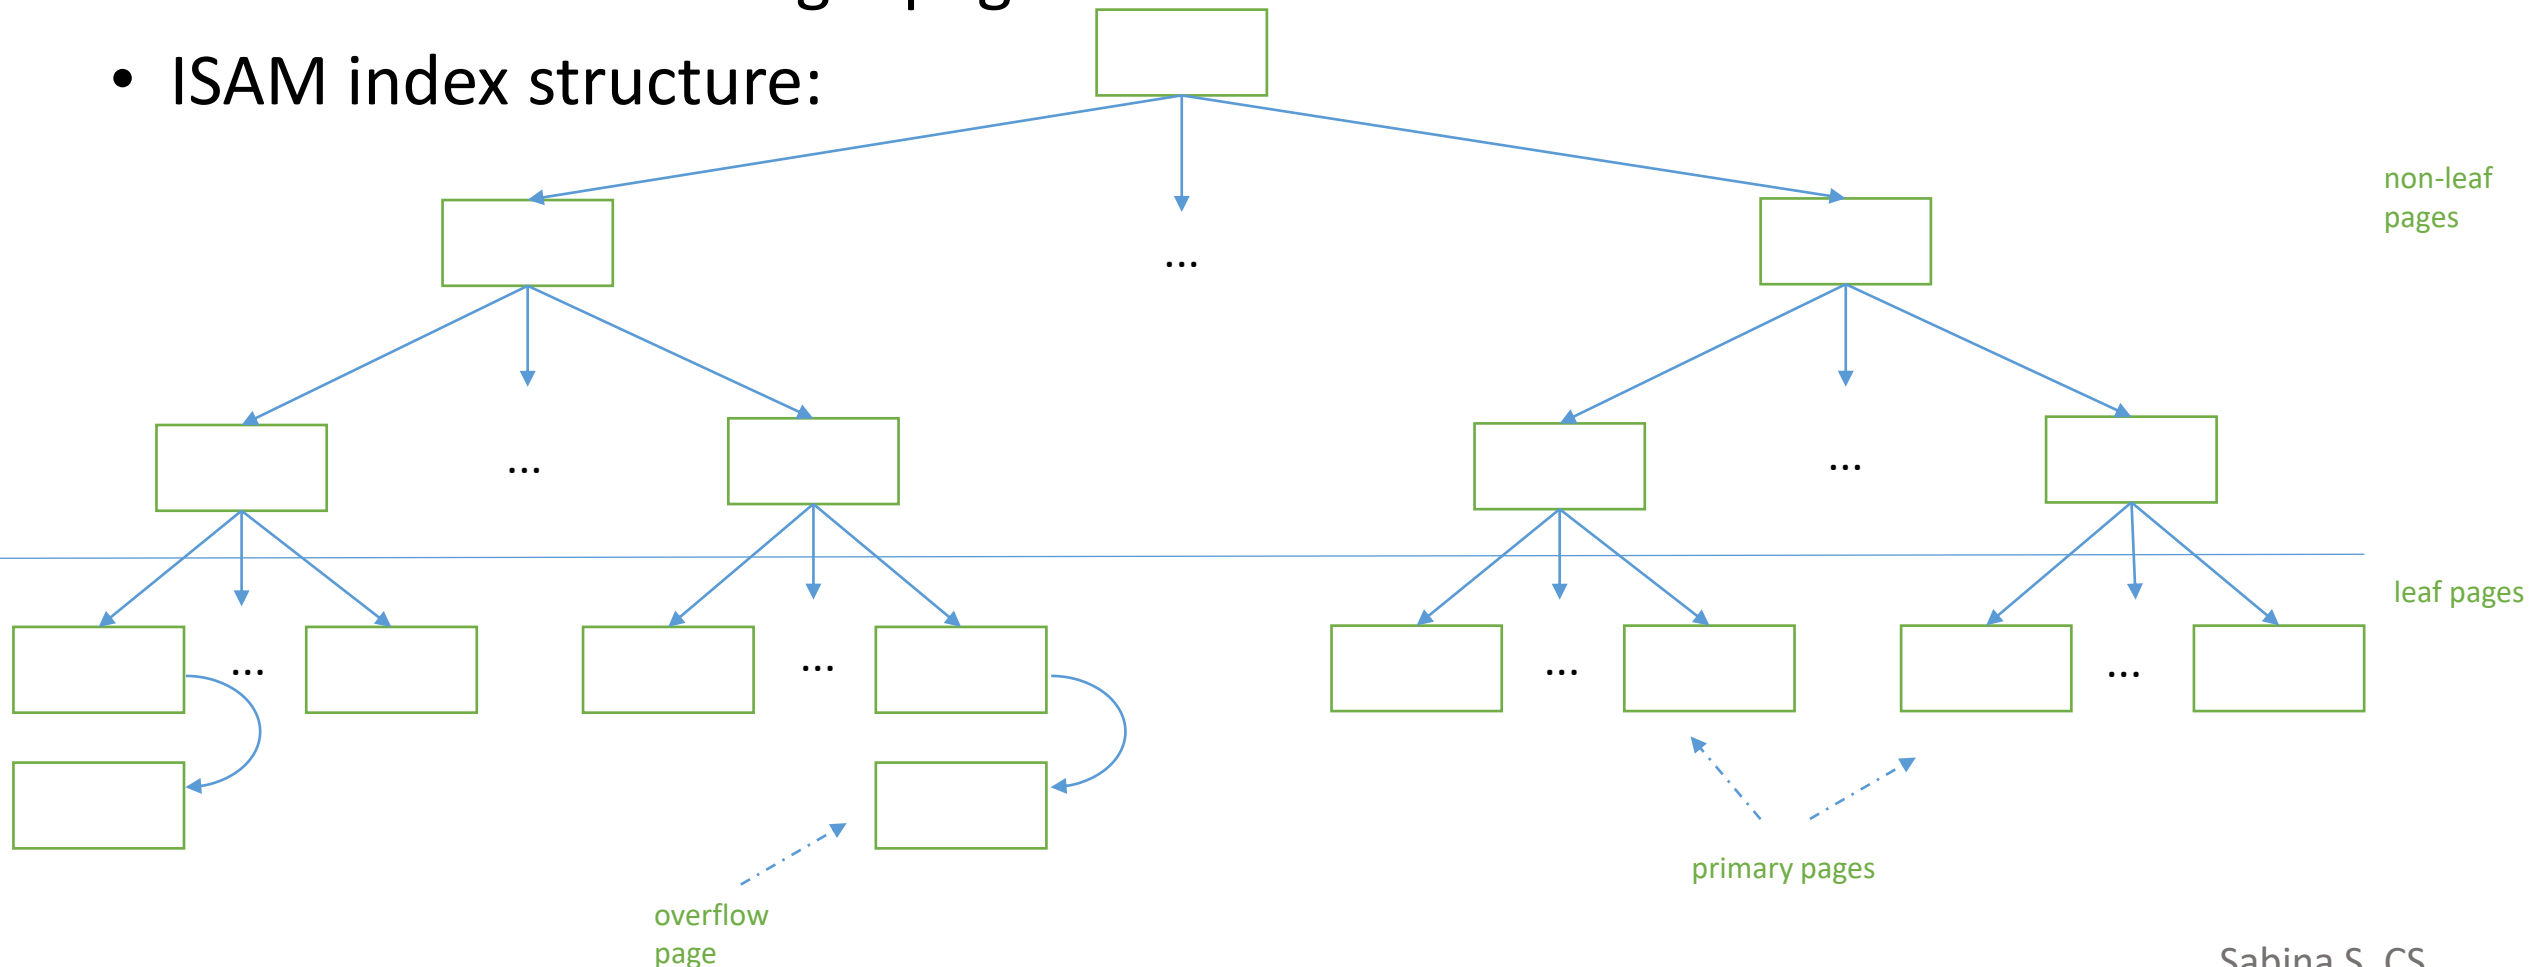
- * Example. Q: Find all phones with *rating* > 9 - range selection query
- data stored in sorted file (records sorted by *rating*) - identify 1st phone using binary search; scan file to get the rest of the phones
- large file => potentially expensive binary search
- create another file with records of the form <1st key on the page, pointer to the page>, sorted on the key (*rating* in the example)



- one-level index structure:



- size of index file - much smaller than size of data file => faster binary search
- index file can still be quite large => further optimization: auxiliary structures are created recursively on top of previously created ones, until one such structure fits on a single page
- ISAM index structure:



- file creation
 - allocate leaf pages - sequentially allocated, sorted on the key
 - allocate non-leaf pages
 - inserts that exceed a page's capacity – allocate overflow pages
- search
 - starts at the root
 - comparisons with the key to find the leaf page
 - cost – disk I/O

data pages

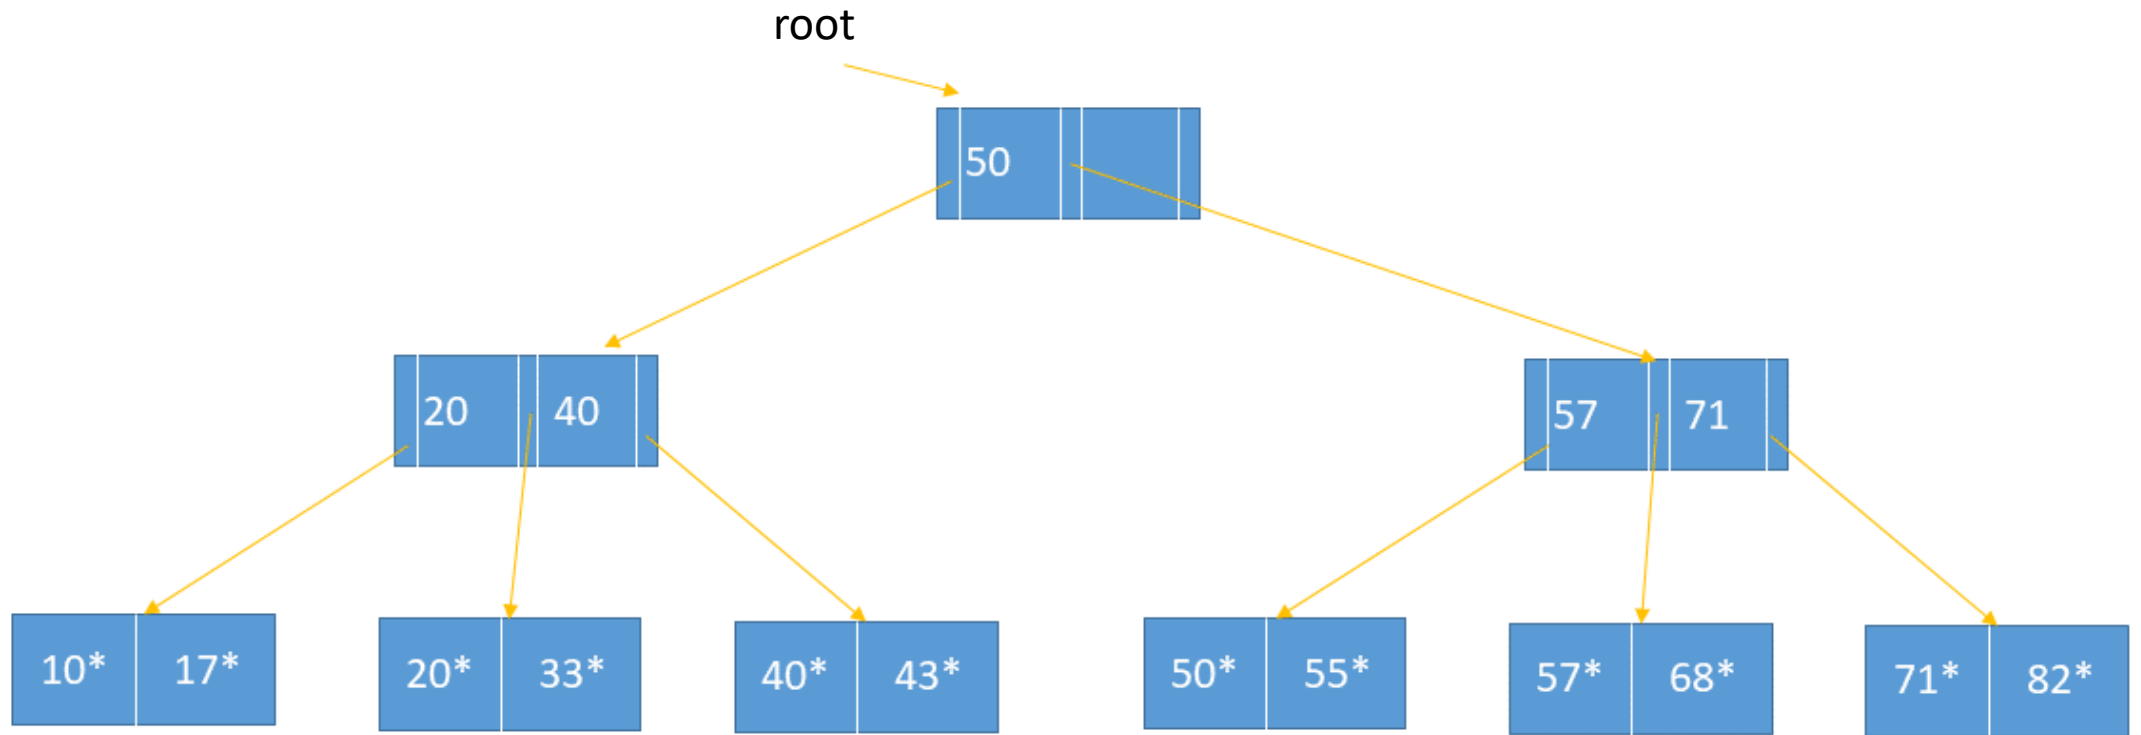
index pages

overflow pages

- insertion
 - find the corresponding leaf page, add the entry
 - if there is no space on the page, add an overflow page
- deletion
 - find the leaf page that contains the entry, remove the entry
 - if an overflow page is emptied, it can be eliminated
- inserts / deletes
 - only leaf pages are affected (static structure)

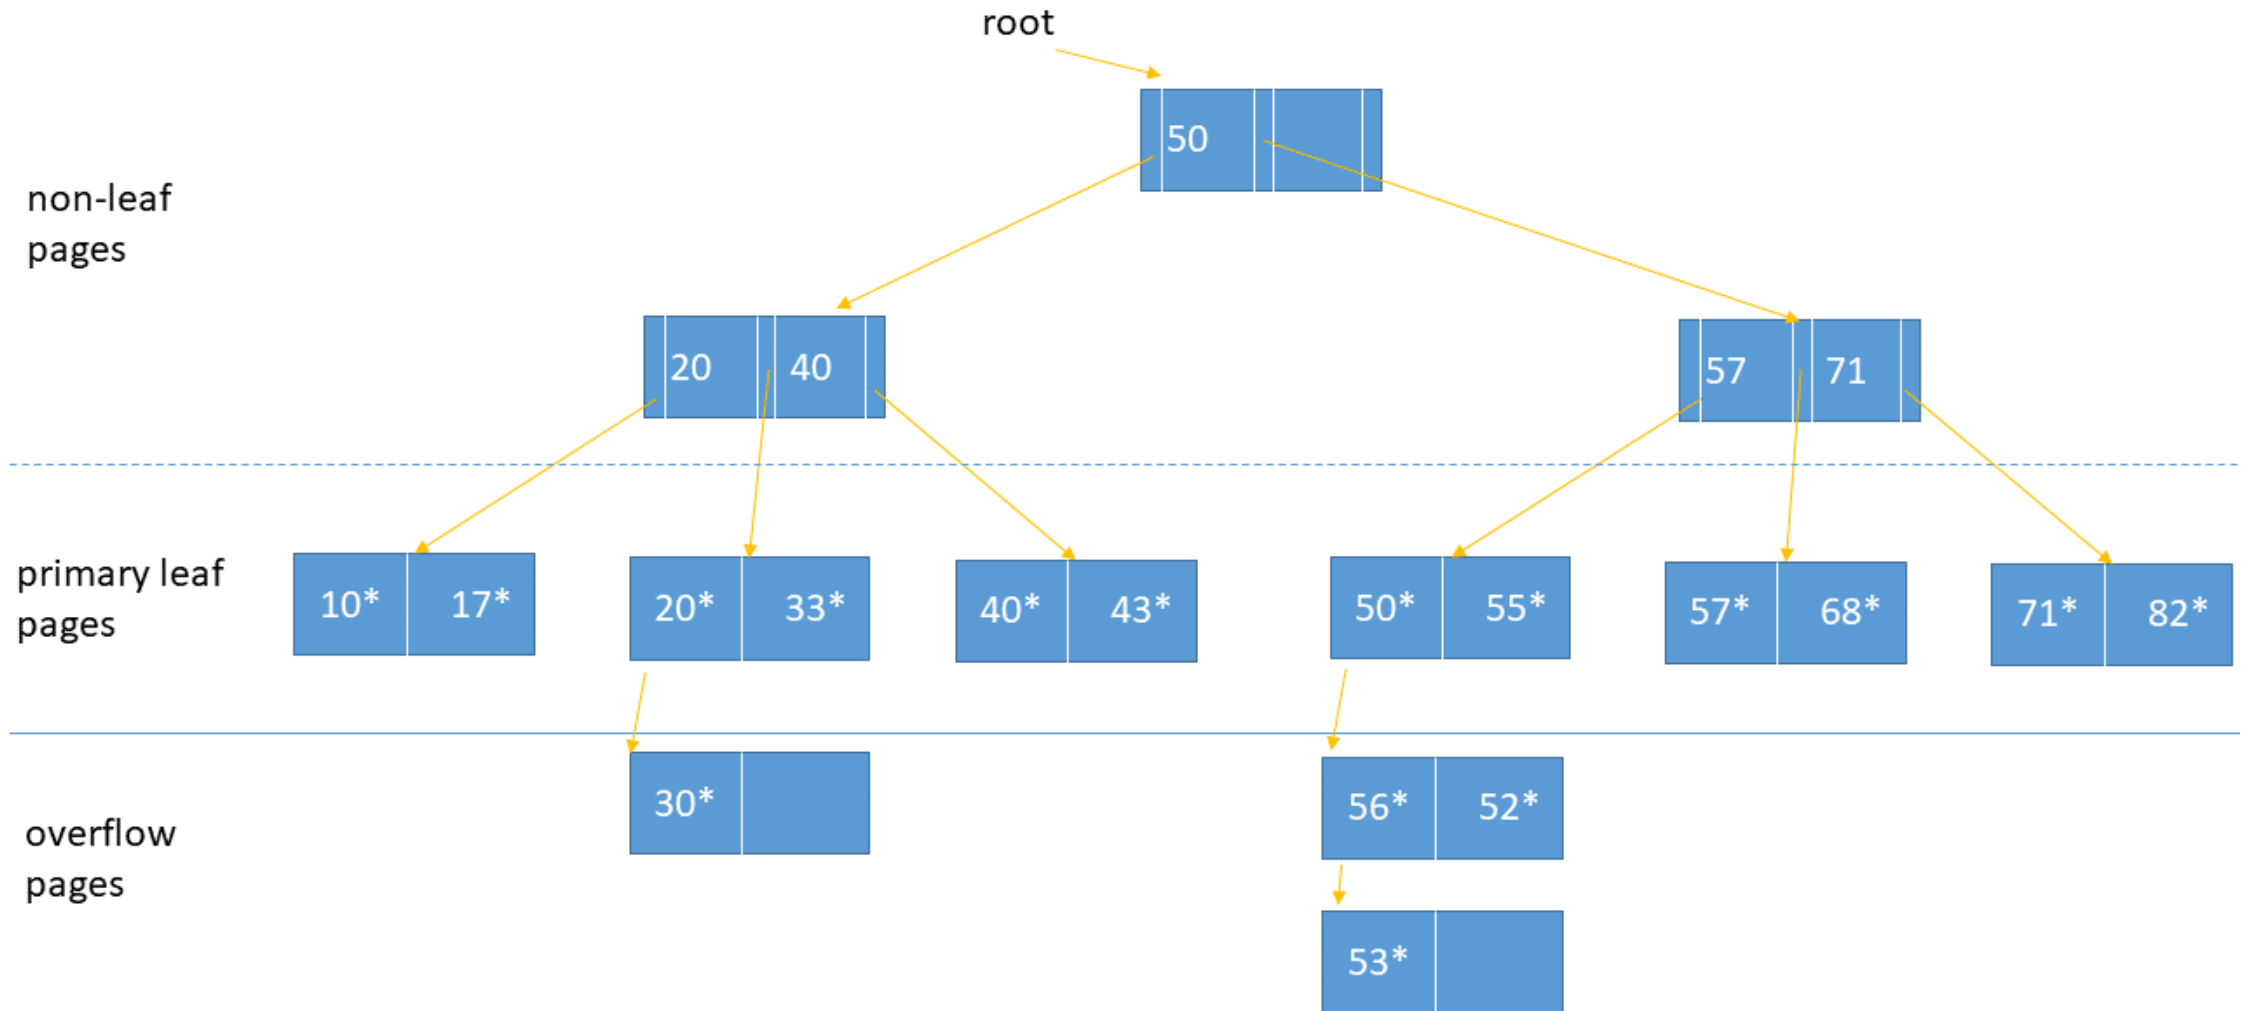
* Example - ISAM tree

- leaf page - 2 entries

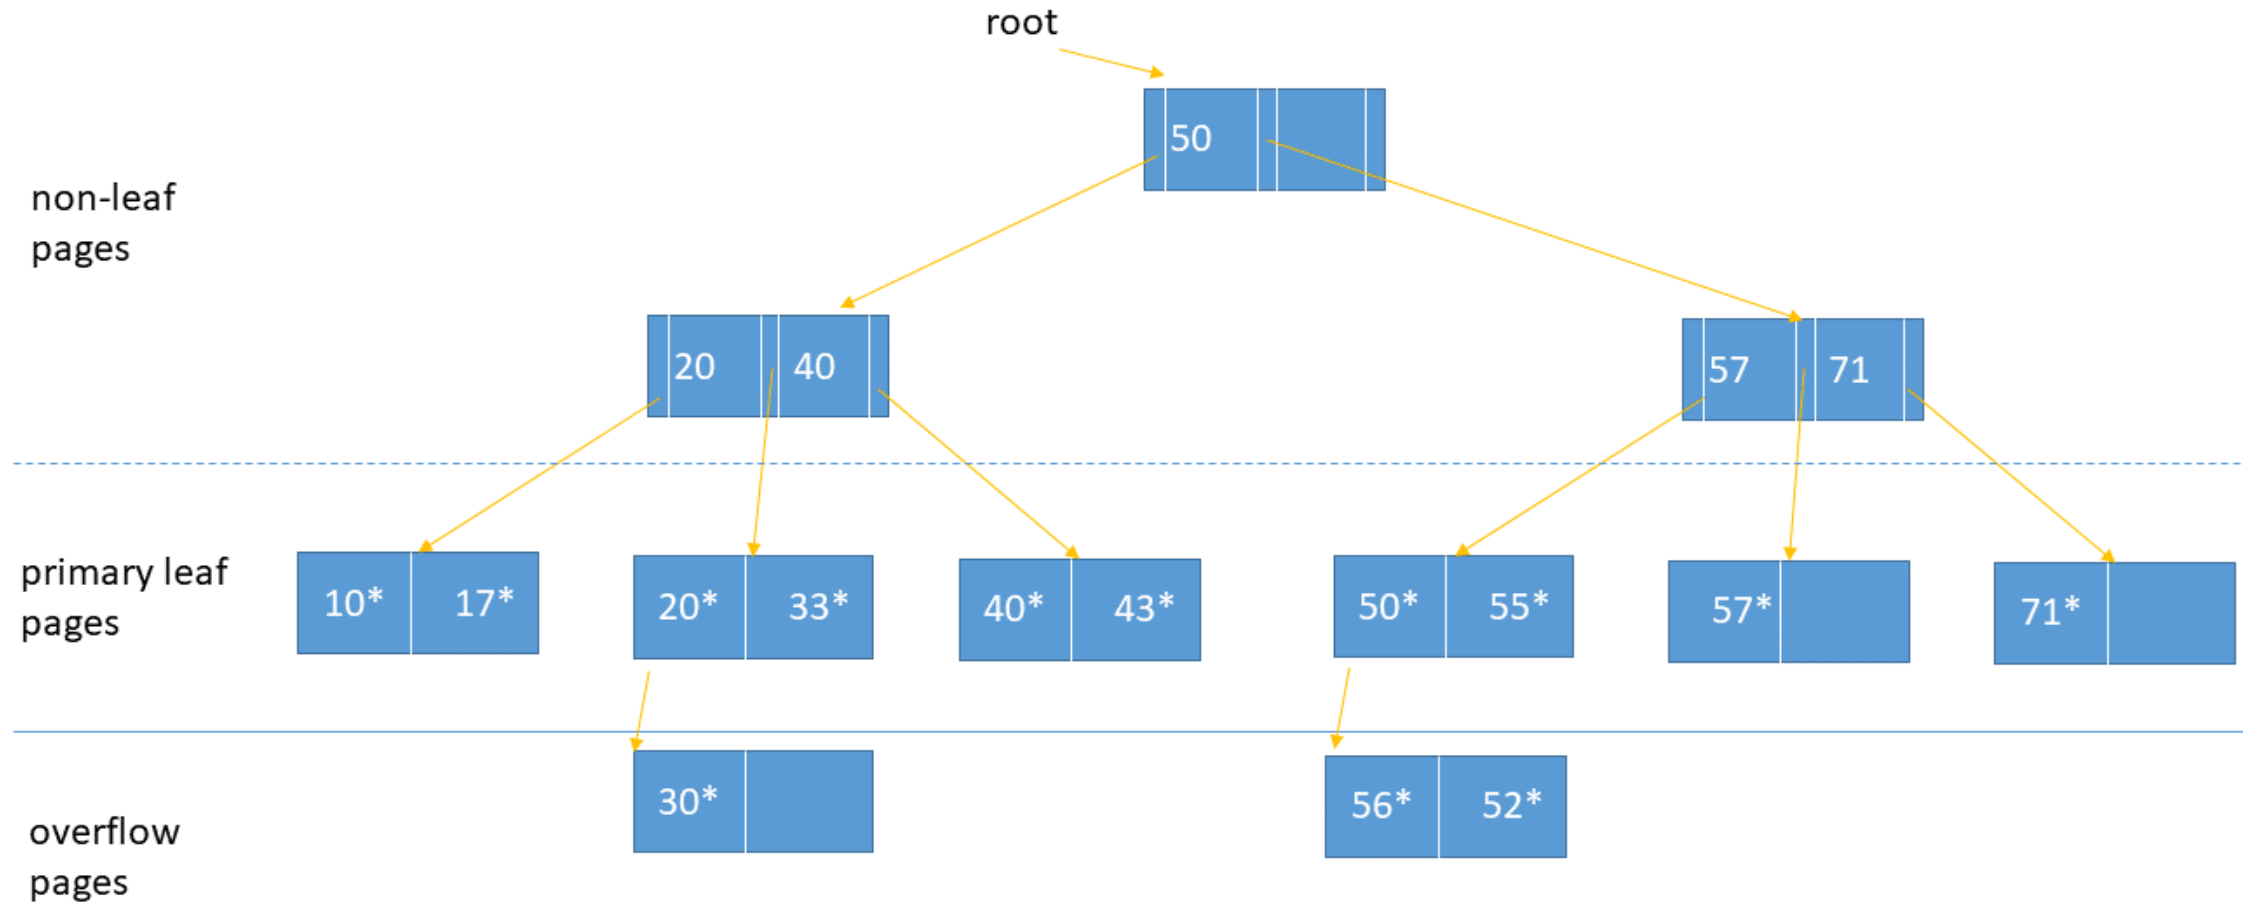


- only key values are shown

- after inserting 30*, 56*, 52*, 53*



- after deleting 53*, 68*, 82*

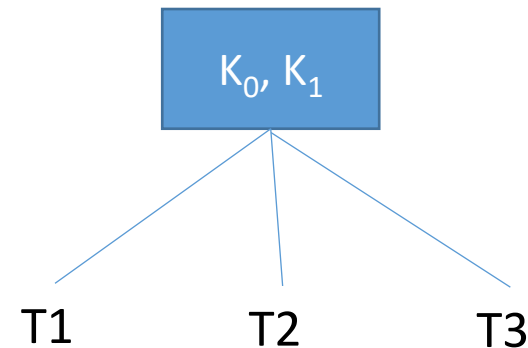
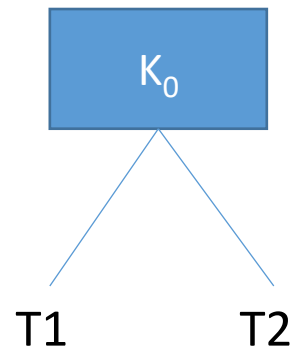


- benefits and drawbacks
 - better concurrent access, since only leaf pages are modified
 - long overflow chains can develop
 - usually not sorted (to optimize inserts)
 - irregular search time if structure not balanced
 - eliminated through deletes / file reorganization
 - when creating the tree - 20% of each page free for future inserts
 - ISAM - suitable when data size / distribution are relatively static

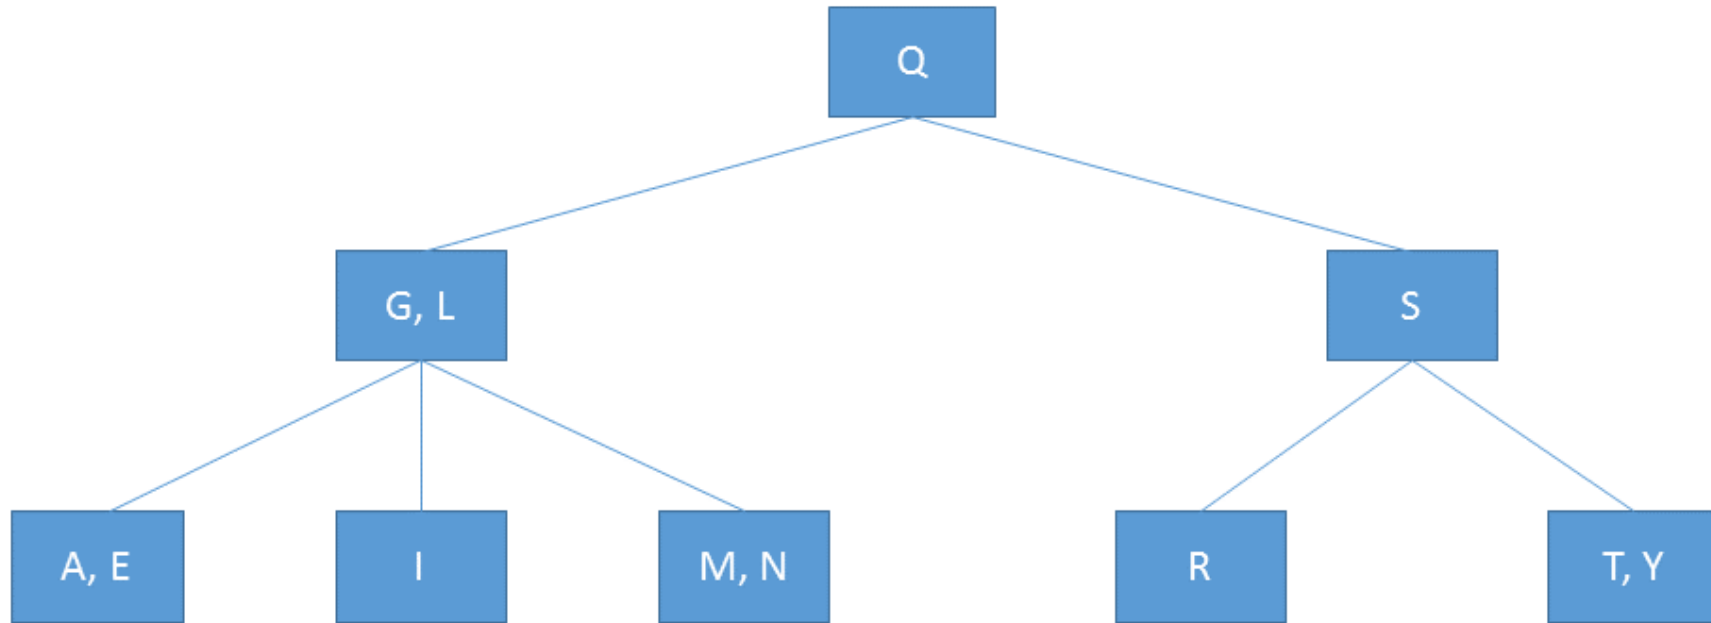
2-3 tree

2-3 tree storing key values (collection of distinct values)

- all the terminal nodes are on the same level
- every node has 1 or 2 key values
 - a non-terminal node with one value K_0 has 2 subtrees: one with values less than K_0 , and one with values greater than K_0
 - a non-terminal node with 2 values K_0 and K_1 , $K_0 < K_1$, has 3 subtrees: one with values less than K_0 , a subtree with values between K_0 and K_1 , and a subtree with values greater than K_1



* Example (key values are letters)



- storing a 2-3 tree
 - 2-3 tree index storing the values of a key
 - tree - key value + address of record (file / DB address of record with corresponding key value)

- 2 options
 1. transform 2-3 tree into a binary tree
 - nodes with 2 values are transformed (see figure below)
 - nodes with 1 value - unchanged



- the structure of a node



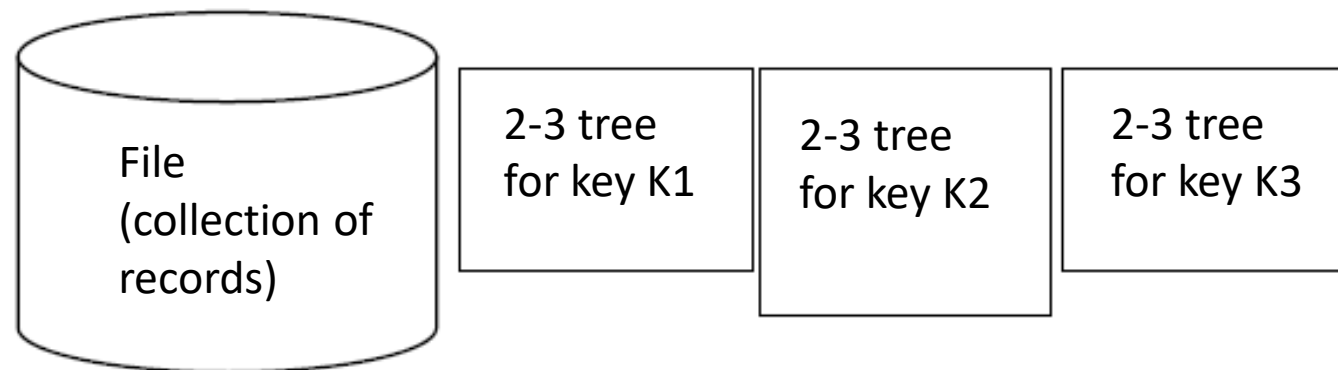
- K - key value
- ADDR - address of the record with the current key value (address in the file)
- PointerL, PointerR - the 2 subtrees' addresses (address in the tree)

- IND - indicator that specifies the type of the link to the right (the 2 possible values can be seen in the previous figure)

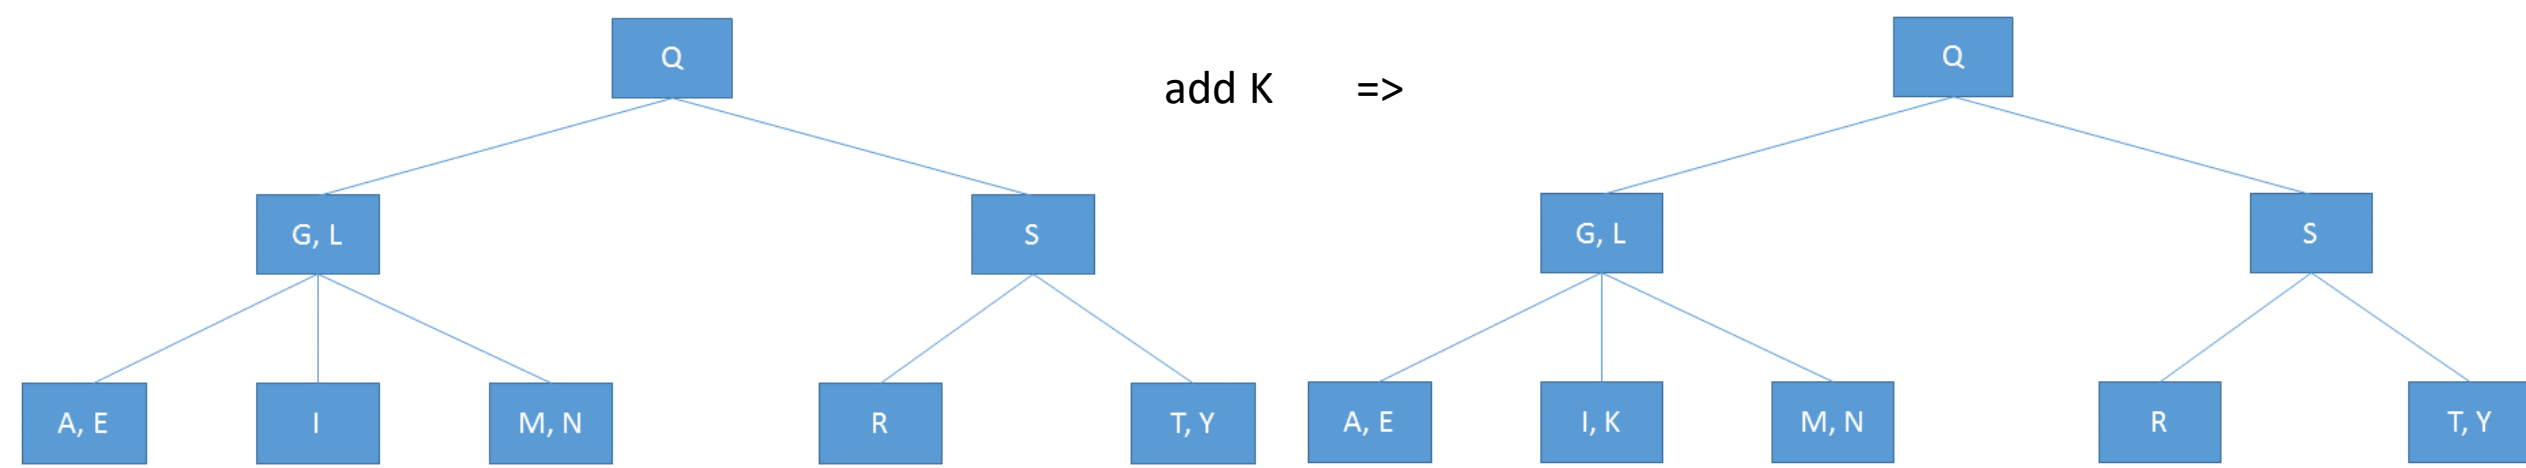
2. the memory area allocated for a node can store 2 values and 3 subtree addresses

NV	K_1	ADDR ₁	K_2	ADDR ₂	Pointer ₁	Pointer ₂	Pointer ₃
----	-------	-------------------	-------	-------------------	----------------------	----------------------	----------------------

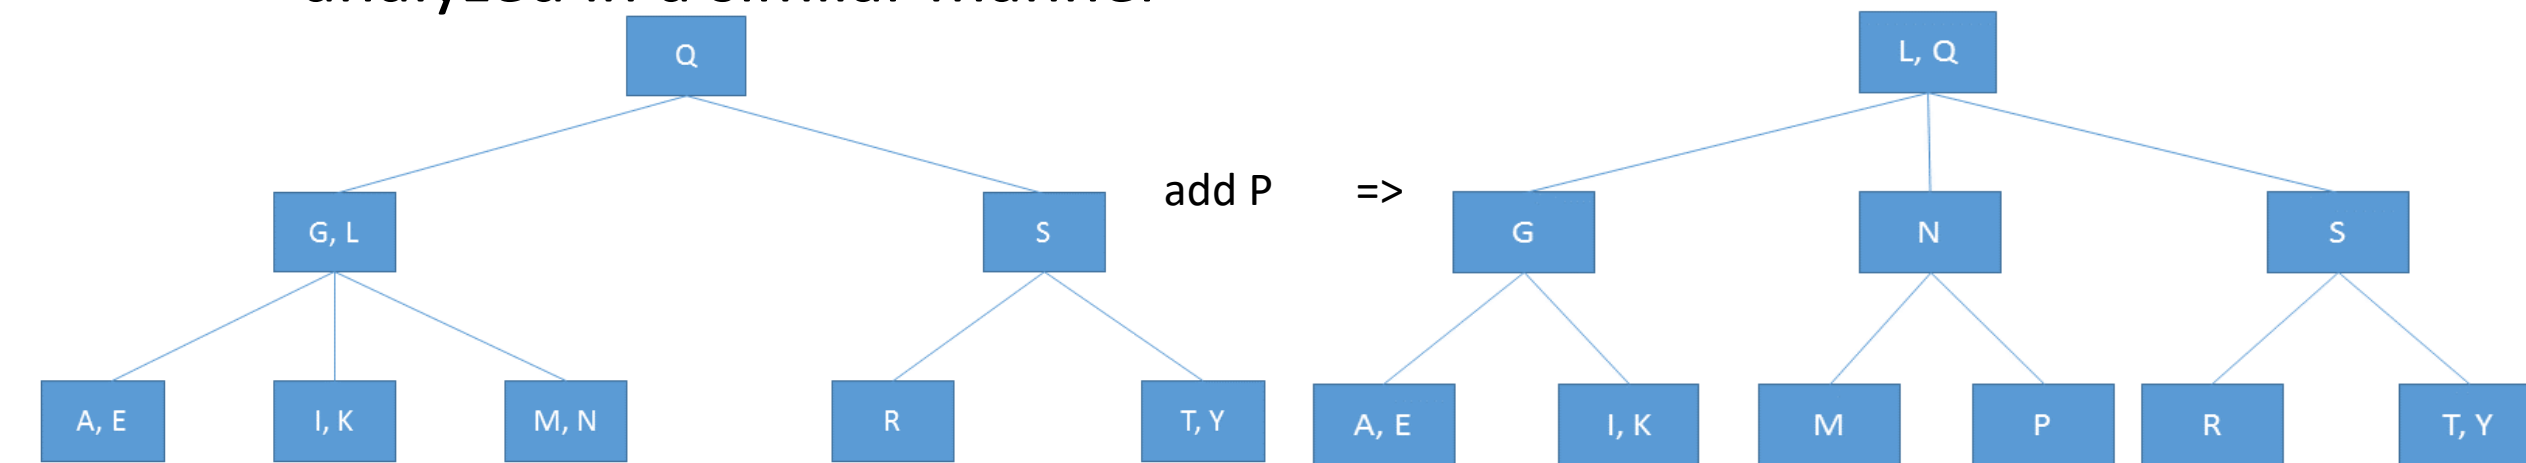
- NV – number of values in the node (1 or 2)
- K_1, K_2 – key values
- ADDR₁, ADDR₂ – the records' addresses (corresponding to K_1 and K_2)
- Pointer₁, Pointer₂, Pointer₃ – the 3 subtrees' addresses
- obs. a file (a relation in a relational DB) can have several associated 2-3 trees (one tree / key)



- operations in a 2-3 tree
 - searching for a record with key value K_0
 - inserting a record - description
 - removing a record - description
 - tree traversal (partial, total)
- add a new value
 - values in the tree must be distinct (the new value should not exist in the tree)
 - perform a test: search for the value in the tree; if the new value can be added, the search ends in a terminal node
 - if the reached terminal node has 1 value, the new value can be stored in the node



- if the reached terminal node has 2 values, the new value is added to the node, the 3 values are sorted, the node is split into 2 nodes: one node will contain the smallest value, the 2nd node - the largest value, and the middle value is attached to the parent node; the parent is then analyzed in a similar manner

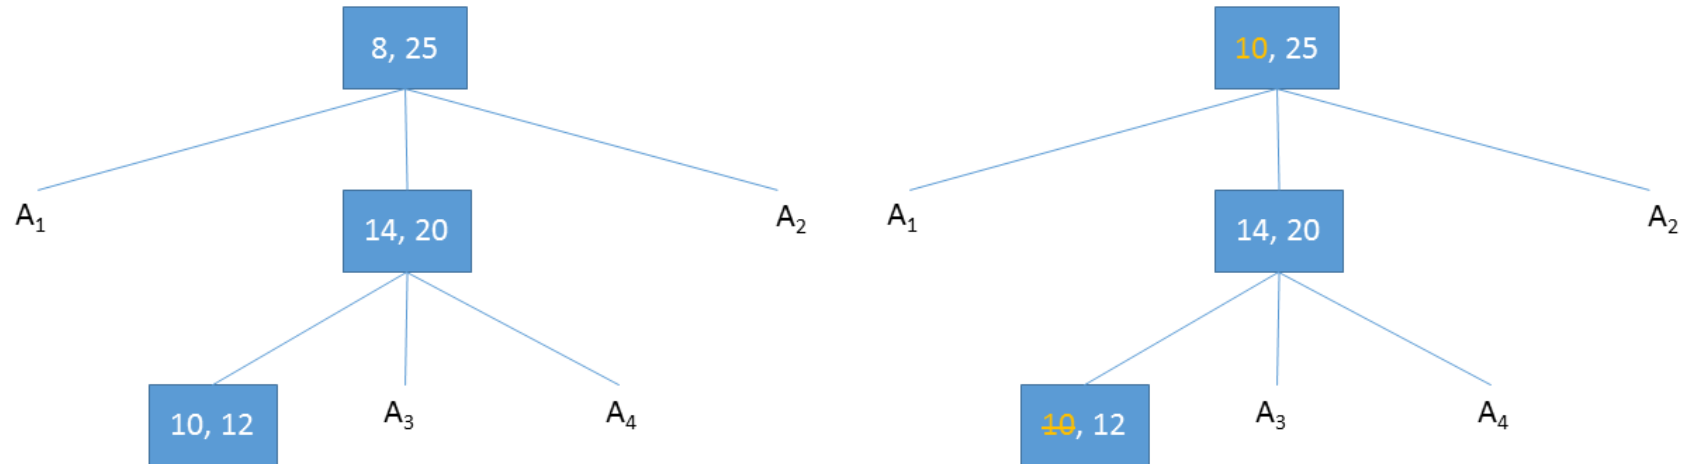


- delete a value K_0

1. search for K_0 ; if K_0 appears in an inner node, change it with a neighbor value K_1 from a terminal node (there is no other value between K_0 and K_1)

- K_1 's previous position (in the terminal node) is eliminated

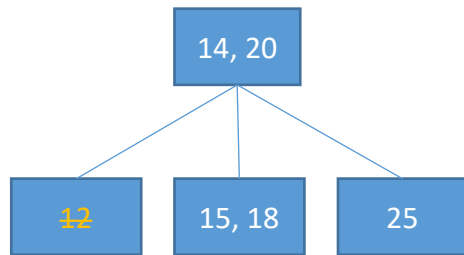
- e.g., remove 8:



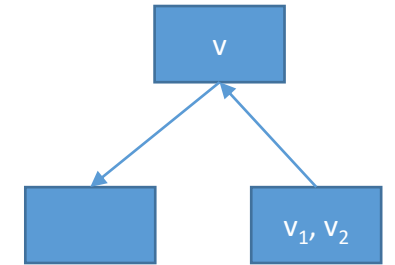
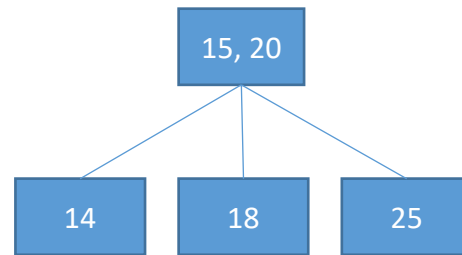
2. perform this step until case a / b occurs

a. if the current node (from which a value is removed) is the root or a node with 1 remaining value, the value is eliminated; the algorithm ends

b. if the delete operation empties the current node, but 2 values exist in one of the sibling nodes (left / right), 1 of the sibling's values is transferred to the parent, 1 of the parent's values is transferred to the current node; the algorithm ends



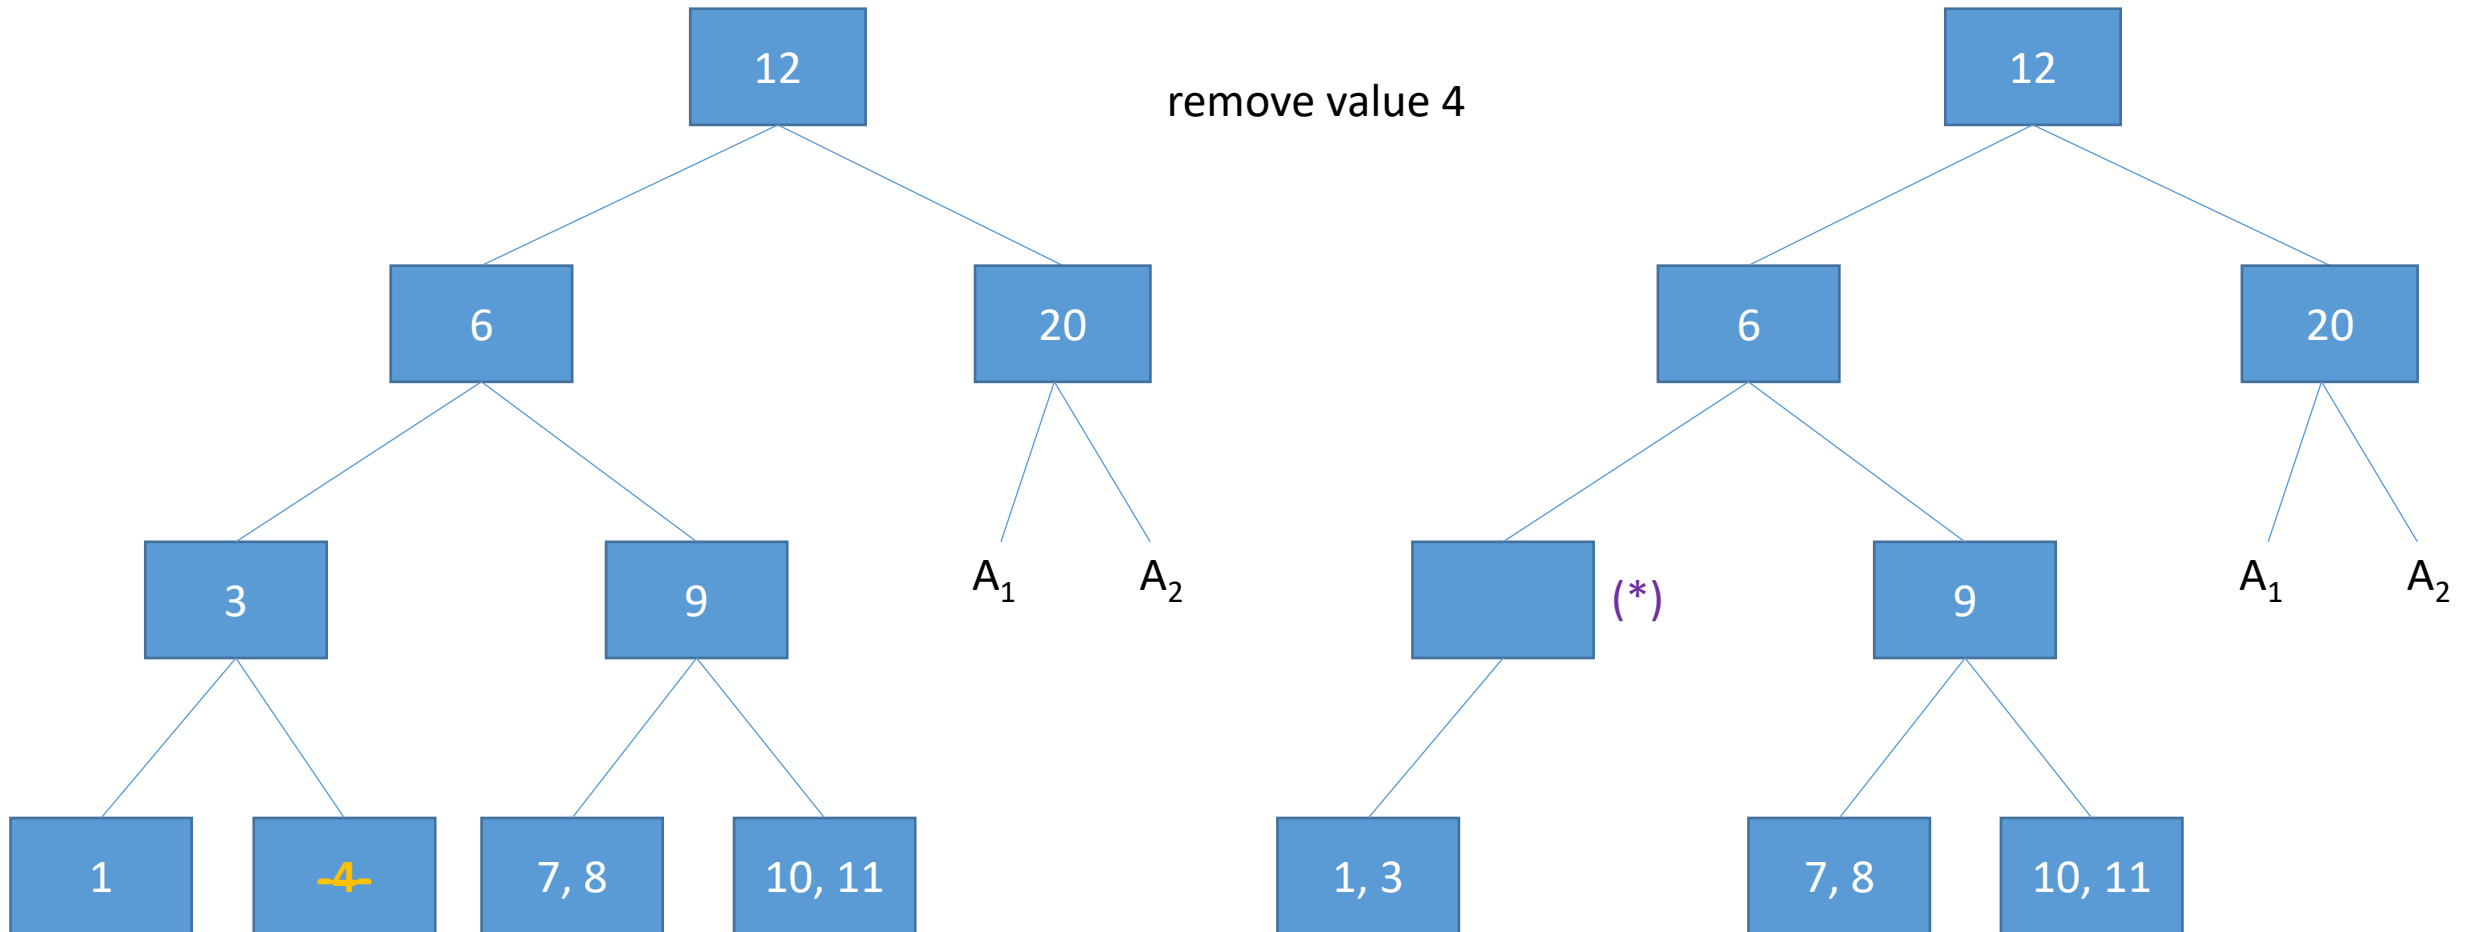
delete value 12

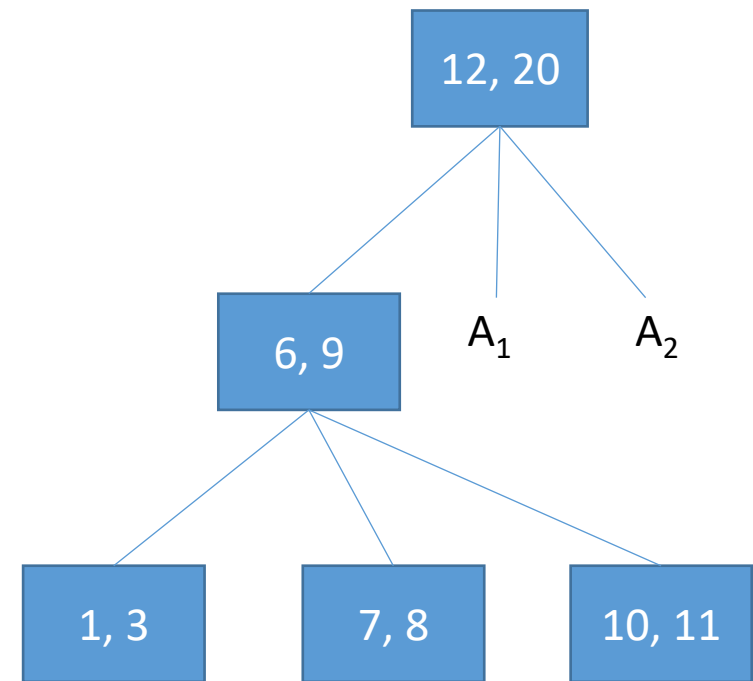
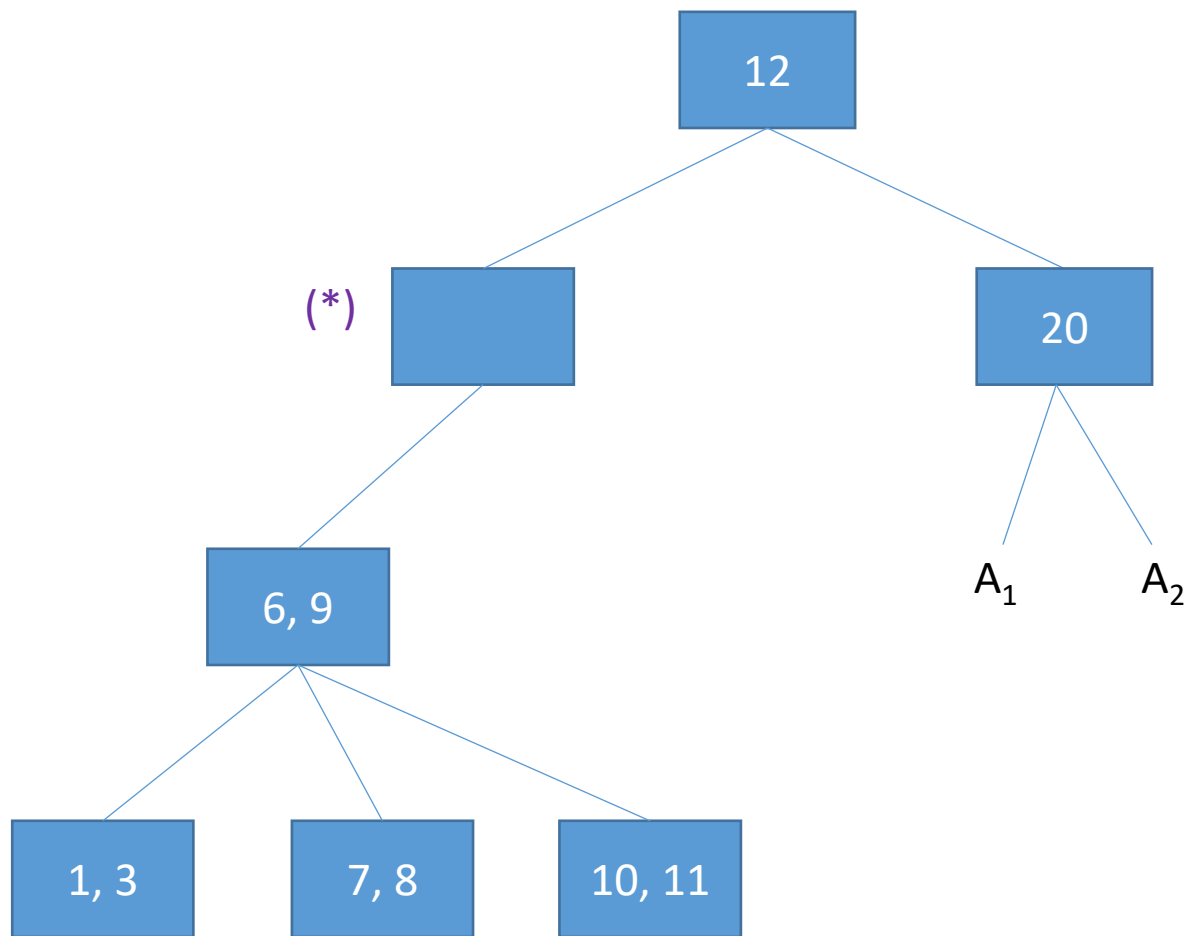


c. if the previous cases do not occur (current node has no values, sibling nodes have 1 value each), then the current node is merged with a sibling and a value from the parent node; case 2 is then analyzed for the parent

- if the root is reached and it has no values, it is eliminated and the current node becomes the root

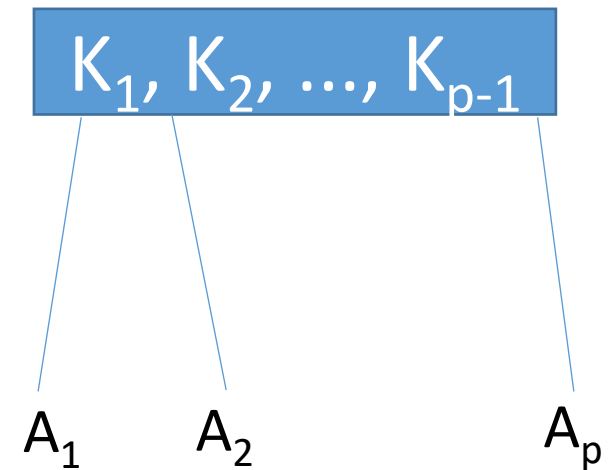
- example: case c for the node marked with (*)





B-tree

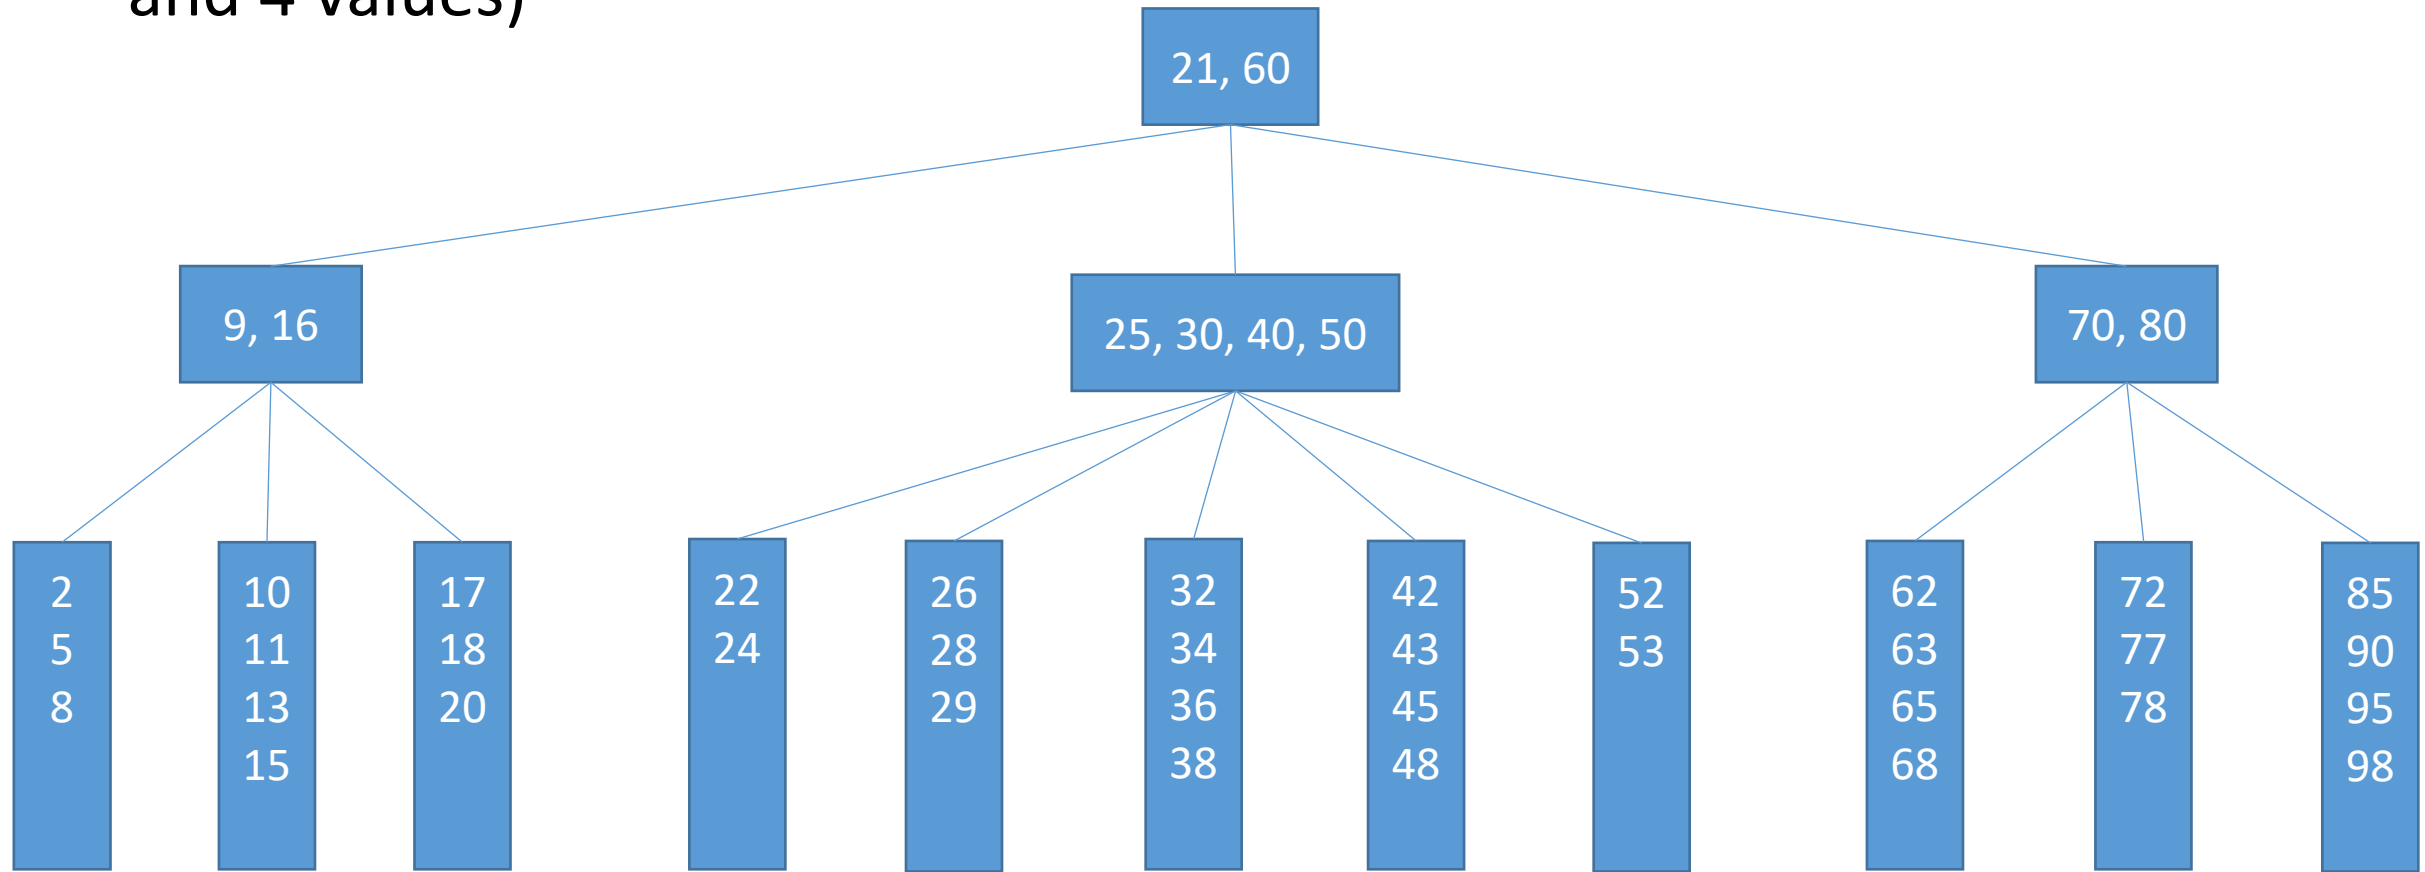
- generalization of 2-3 trees
- B-tree of order m
 1. if the root is not a terminal, it has at least 2 subtrees
 2. all terminal nodes – same level
 3. every non-terminal node – at most m subtrees
 4. a node with p subtrees has p-1 ordered values (ascending order): $K_1 < K_2 < \dots < K_{p-1}$
 - A_1 : values less than K_1
 - A_i : values between K_{i-1} and K_i , $i=2, \dots, p-1$
 - A_p : values greater than K_{p-1}



5. every non-terminal node – at least $\left\lceil \frac{m}{2} \right\rceil$ subtrees
- obs. limits on number of subtrees (and values) / node result from the manner in which inserts / deletes are performed so that the second requirement in the definition is met

* Example - B-tree of order 5

- non-terminal, non-root node – at most 5, at least 3 subtrees (between 2 and 4 values)



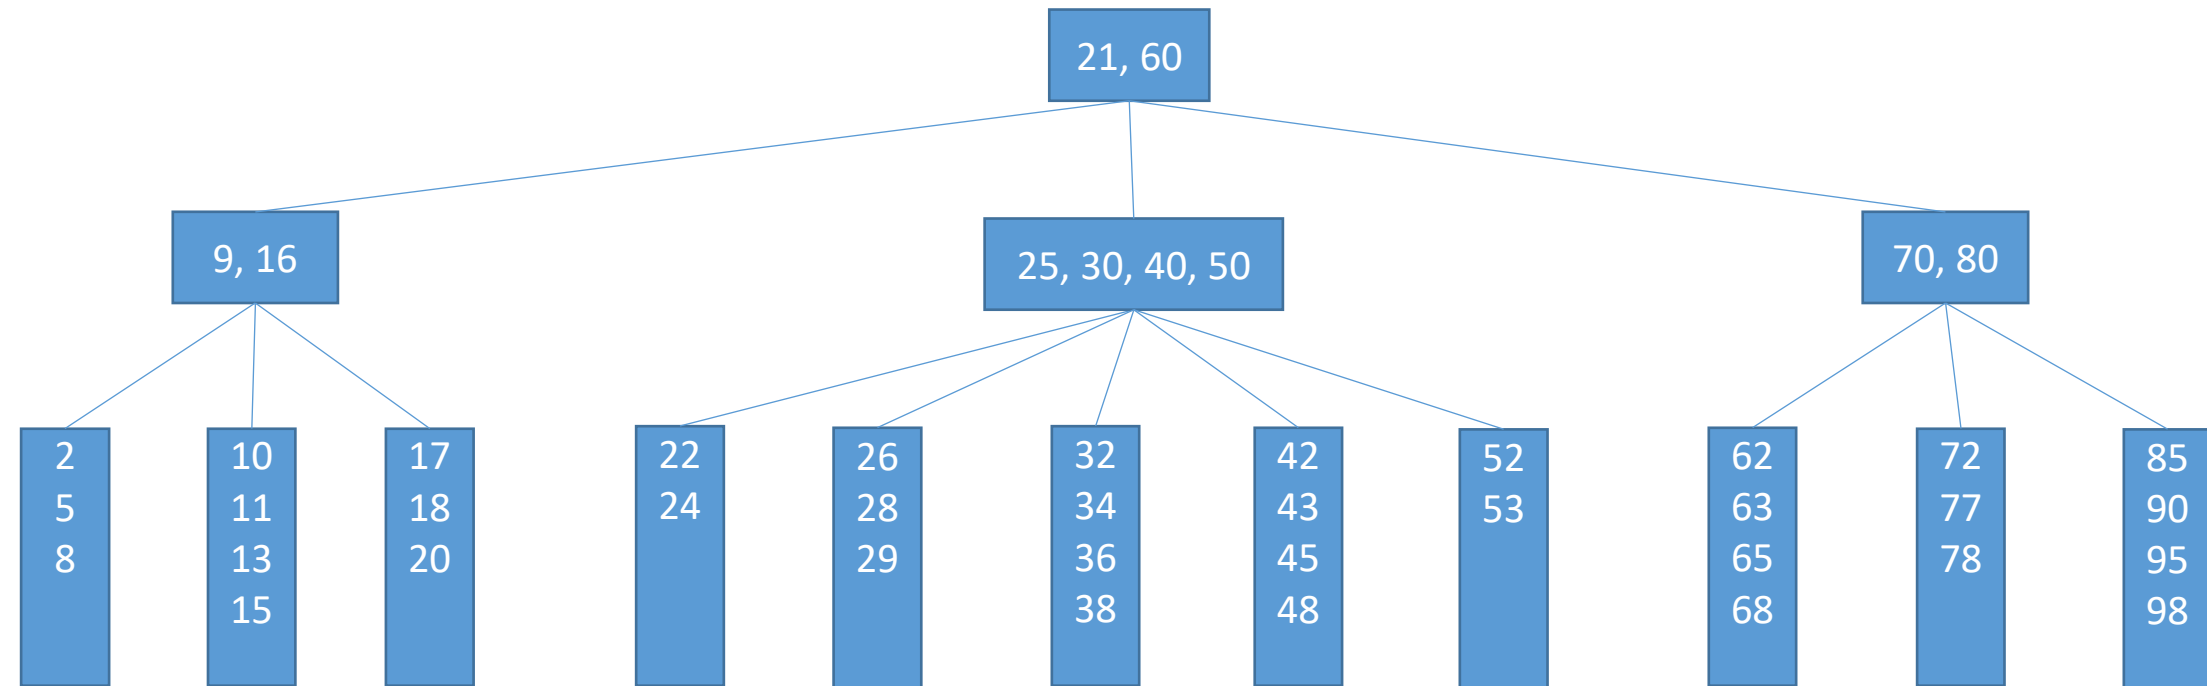
- B-tree of order m
 - storing the values of a key (a database index)
 - tree
 - key value + address of record
- 1. transformed into a binary tree
 - 2-3 tree method
- 2. the memory area allocated for a node can store the maximum number of values and subtree addresses

NV	K_1	$ADDR_1$...	K_{m-1}	$ADDR_{m-1}$	$Pointer_1$...	$Pointer_m$
----	-------	----------	-----	-----------	--------------	-------------	-----	-------------

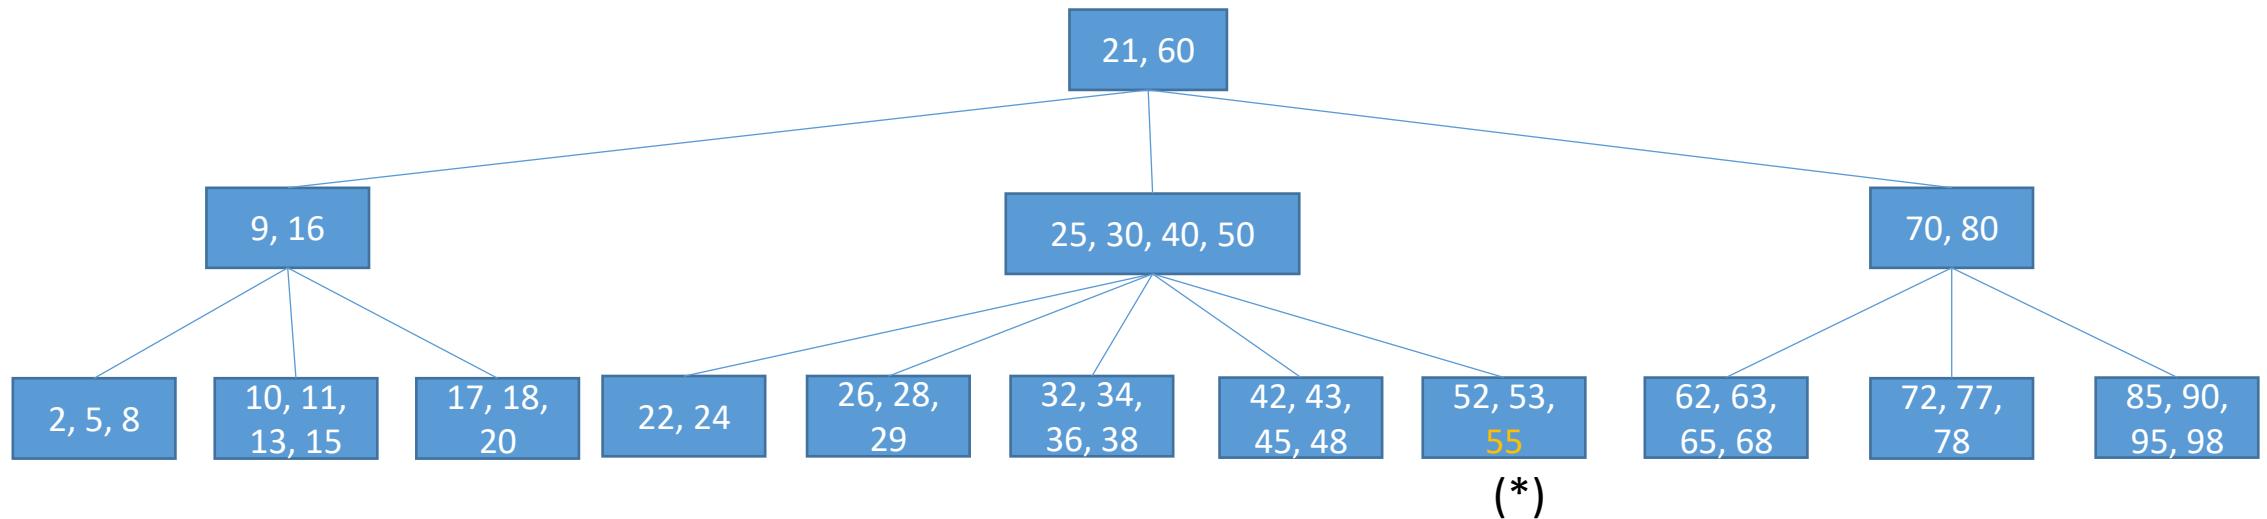
- NV - number of values in the node
- K_1, \dots, K_{m-1} - key values
- $ADDR_1, \dots, ADDR_{m-1}$ - the records' addresses (corresponding to the key's values)
- $Pointer_1, \dots, Pointer_m$ – subtree addresses

- B-tree of order m
 - useful operations in a B-tree
 - searching for a value
 - adding a value - description
 - removing a value- description
 - tree traversal (partial, total)

- B-tree of order m
 - adding a new value
 1. values in the tree must be distinct (the new value should not exist in the tree); perform a test (search for the value in the tree)
 - if the new value can be added, the search ends in a terminal node
 2. if the reached terminal node has less than $m-1$ values, the new value can be stored in the node, e.g., 55 is added to the tree below:

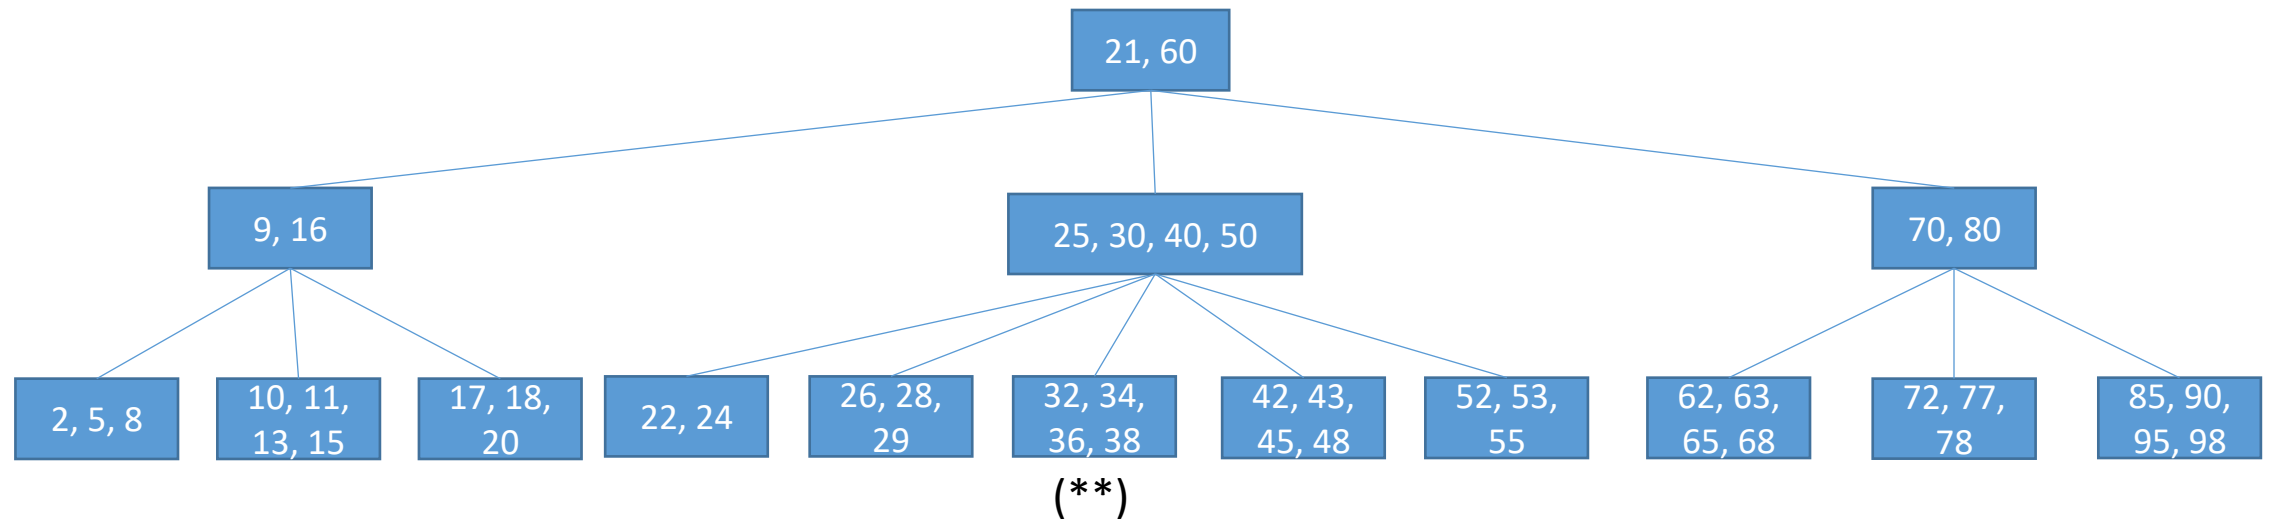


- B-tree of order m
 - adding a new value
 - the resulting tree is shown below:



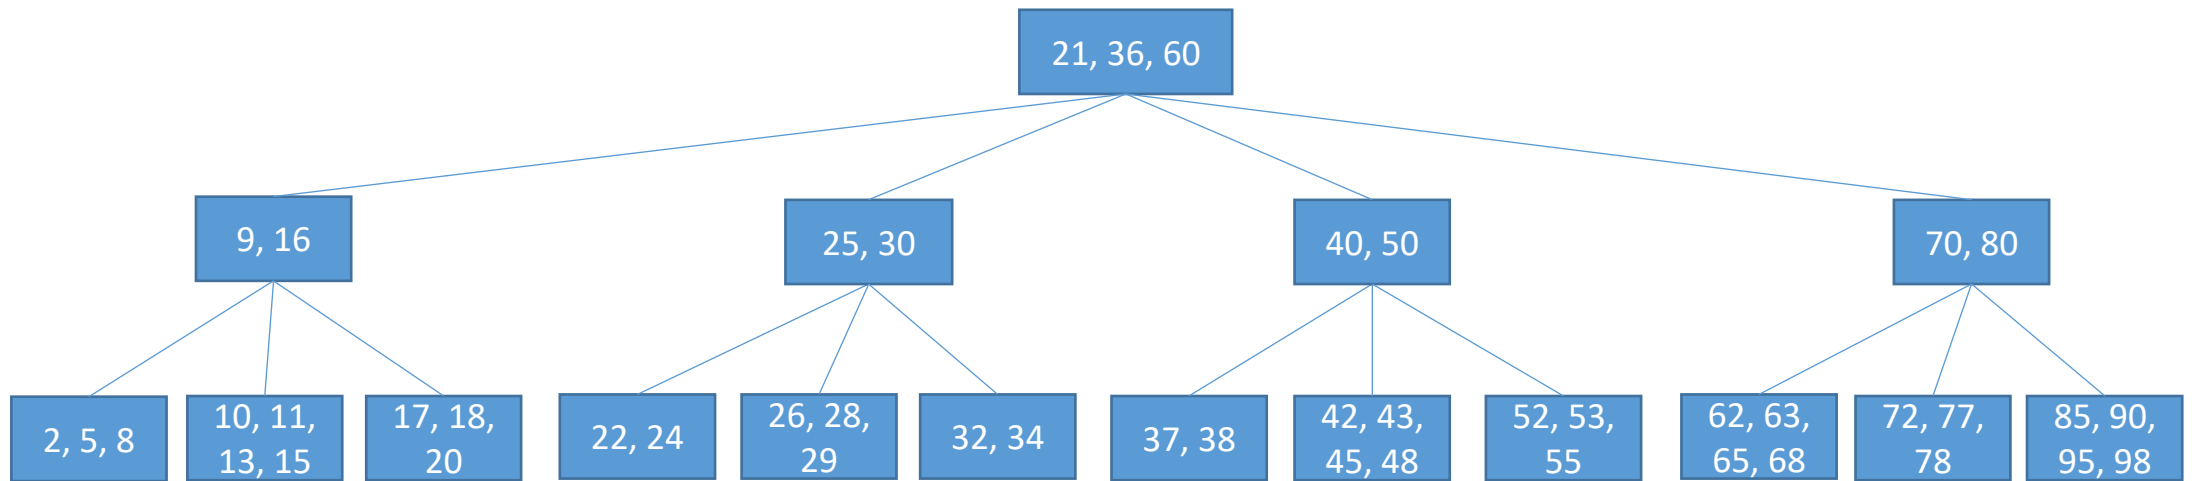
- 55 belongs to the node marked with (*), which can store at most 4 values

- B-tree of order m
 - adding a new value
 3. if the terminal node already has $m-1$ values, the new value is attached to the node, the m values are sorted, the node is split into 2 nodes, and the middle value (median) is attached to the parent node; the parent is then analyzed in a similar manner
 - e.g., add 37 to the tree below

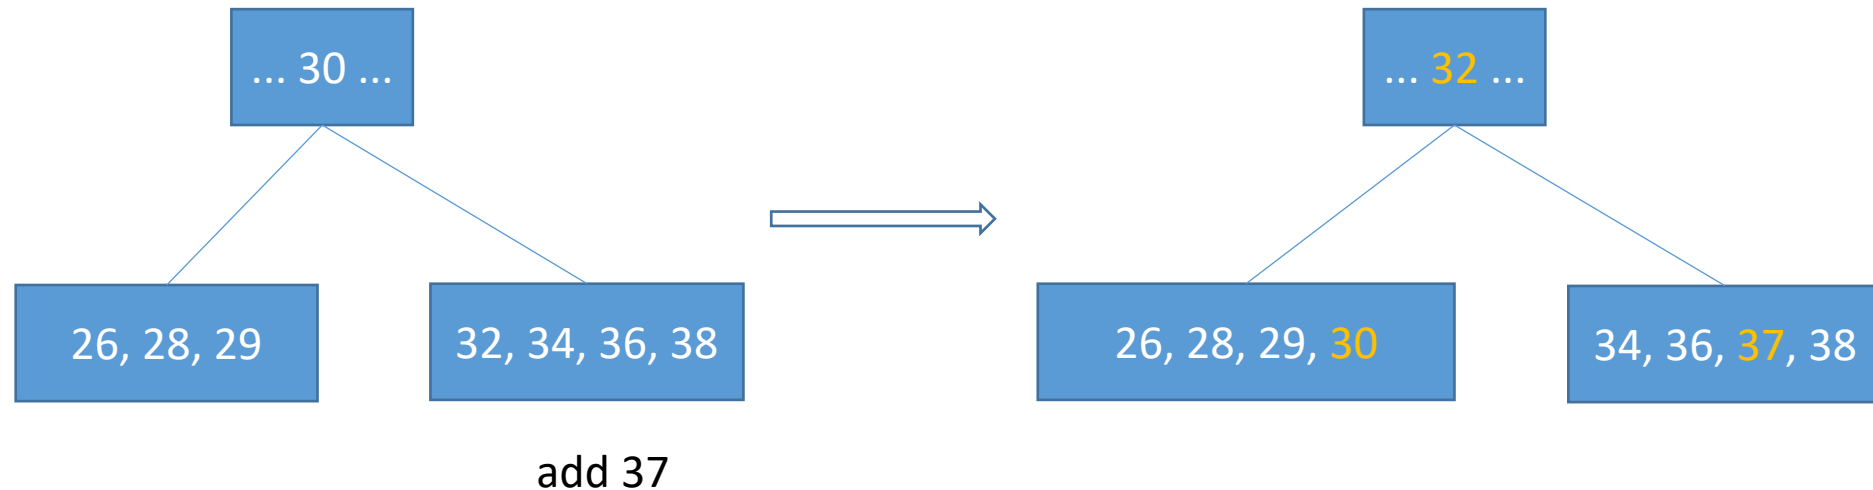


- the node marked with (**) should contain values 32, 34, 36, 37, 38

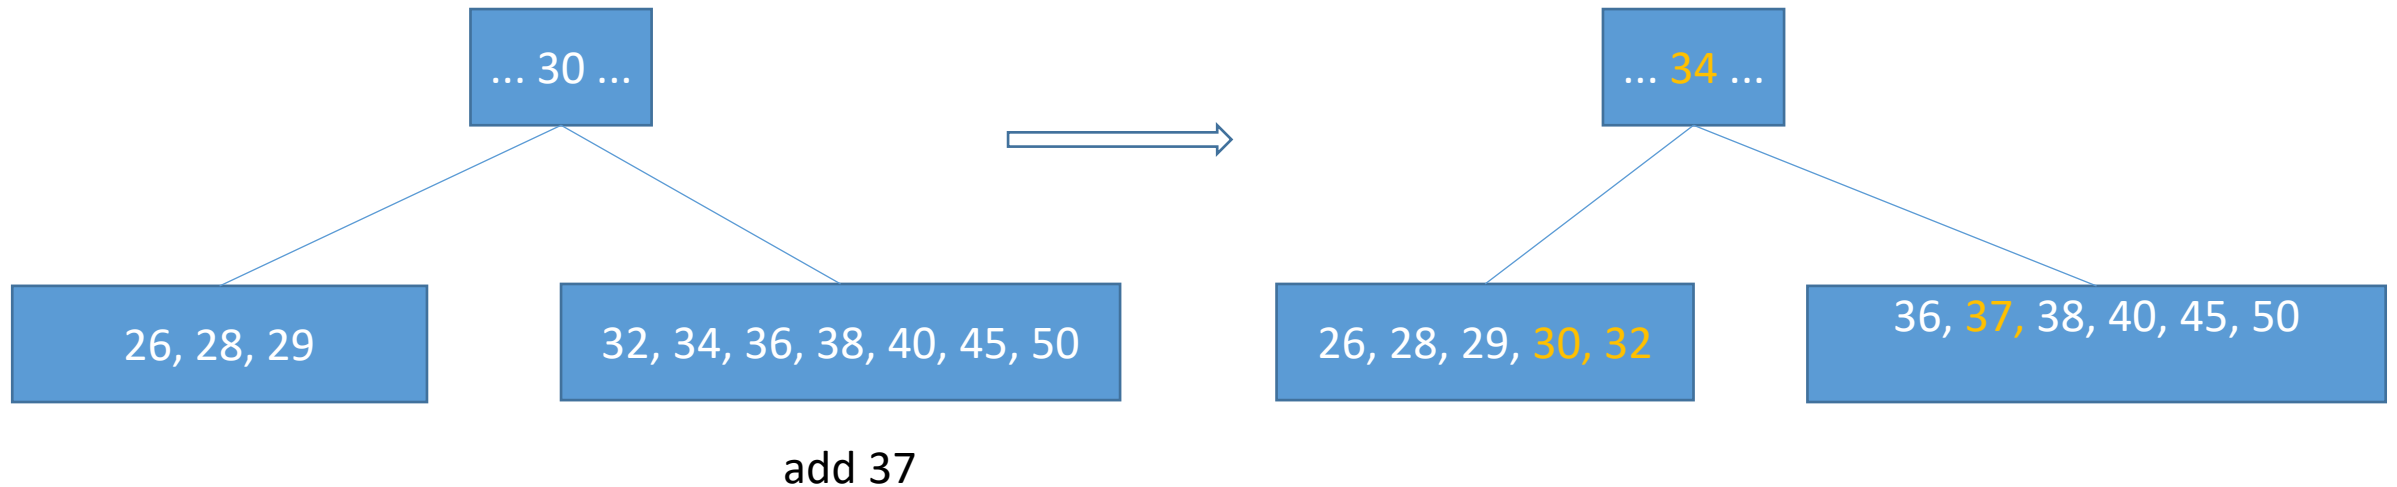
- B-tree of order m
 - adding a new value
 - since the node's capacity is exceeded, it is split into nodes 32, 34, and 37, 38, and 36 is attached to the parent node (with values 25, 30, 40, 50)
 - in turn, the parent must be split into 2 nodes (values 25, 30, and 40, 50), and 36 is attached to its parent



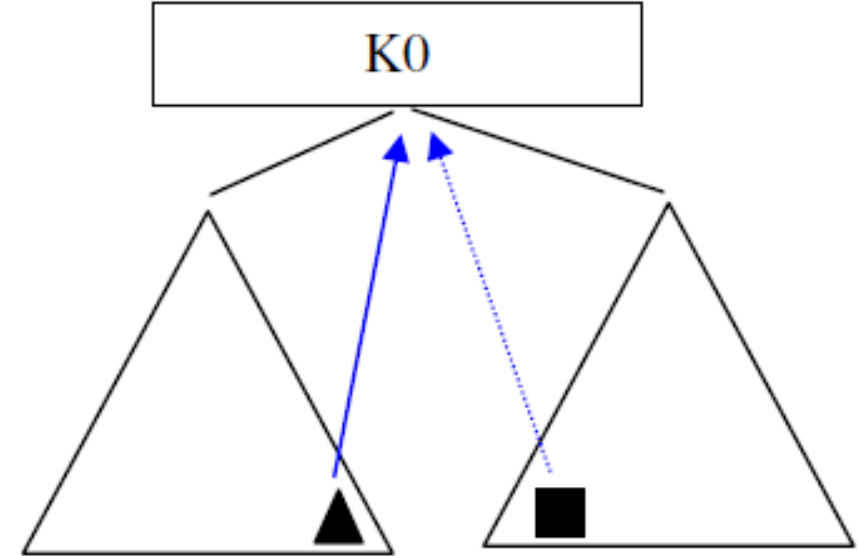
- B-tree of order m
 - adding a new value
 - optimizations
 - before performing a split - analyze whether one or more values can be transferred from the current node (with $m-1$ values) to a sibling node
 - e.g., B-tree of order 5 (non-terminal node - between 2 and 4 values, i.e., between 3 and 5 subtrees):



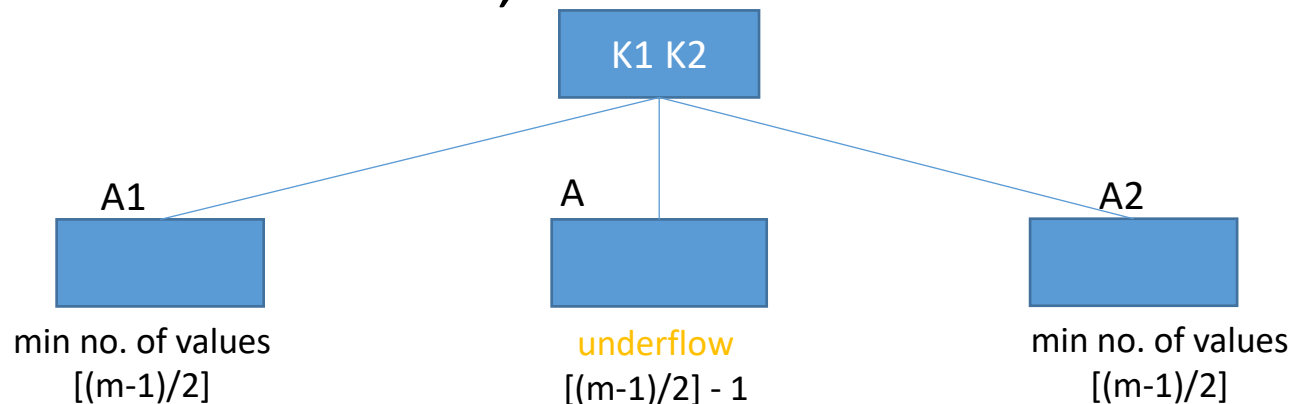
- B-tree of order m
 - adding a new value
 - optimizations
 - e.g., B-tree of order 8 (non-terminal node - between 3 and 7 values, i.e., between 4 and 8 subtrees):



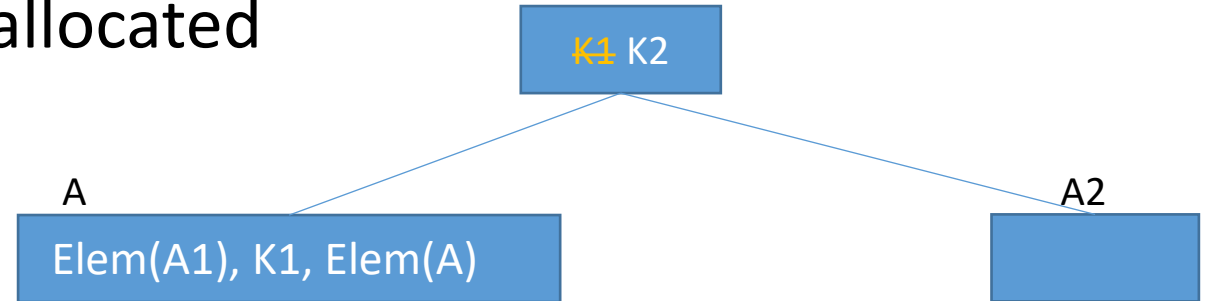
- B-tree of order m
 - removing a value
 - a node can have at most m subtrees, i.e., a maximum of m-1 values, and at least $\lceil \frac{m}{2} \rceil$ subtrees, i.e., at least $\lceil \frac{m}{2} \rceil - 1 = \lceil \frac{m-1}{2} \rceil$ values
 - when eliminating a value from a node, an underflow can occur (the node can end up with less values than the required minimum)
 - eliminate value K_0
 1. search for K_0 ; if it doesn't exist, the algorithm ends
 2. if K_0 is found in a non-terminal node (like in the figure on the right), K_0 is replaced with a *neighbor value* from a terminal node (this value can be chosen between 2 values from the trees separated by K_0)



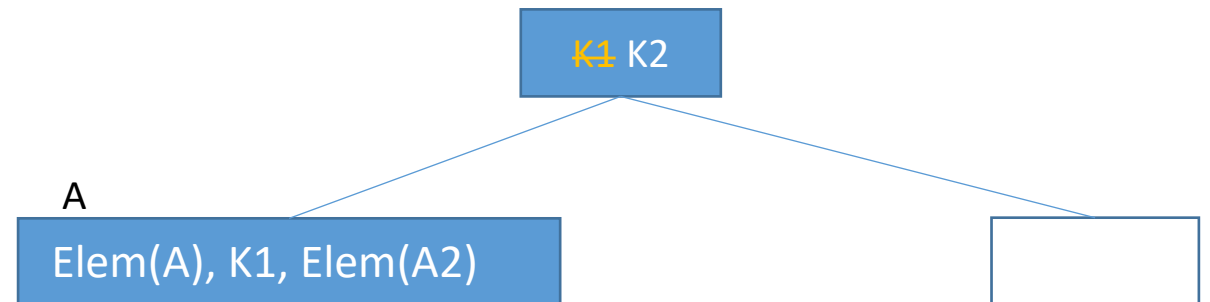
- B-tree of order m
 - removing a value
 3. perform this step until case a / b occurs
 - a. if the current node (from which a value is removed) is the root or underflow doesn't occur, the value is eliminated; the algorithm ends
 - b. if the delete operation causes an underflow in the current node (A), but one of the sibling nodes (left / right - B) has at least 1 extra value, values are transferred between A and B via the parent node; the algorithm ends
 - c. if there is an underflow in A, and sibling nodes A1 and A2 have the minimum number of values, nodes must be concatenated:



- B-tree of order m
 - removing a value
 - if A1 exists, A1 is merged with A and value K1 (separating A1 from A); the node at address A1 is deallocated



- if there is no A1 (A is the first subtree for its parent), A is merged with A2 and K1 (separating A from A2); the node at address A2 is deallocated



- case 3 is then analyzed for the parent node
- if the root is reached and has no values, it is removed and the current node becomes the root

- B-tree of order m
 - obs. a block stores a node from a B-tree
- e.g.:
 - key size: 10b
 - record address / node address: 10b
 - NV value (number of values in the node): 2b
 - block size: 1024b (10b for the header)
- then: $2 + (m-1) * (10+10) + m * 10 = 1024 - 10 \Rightarrow m = 34$
- if the size of a block is 2048b and the other values are unchanged, then the order of the tree is $m = 68$, i.e., a node can have between 33 and 67 values

- B-tree of order m
- the maximum number of required blocks (from the file that stores the B-tree) when searching for a value - the maximum number of levels in the tree; for $m=68$, if the number of values is 1.000.000, then:
 - the root node (on level 0) contains at least 1 value (2 subtrees)
 - on the next level (level 1) - at least 2 nodes * 33 values/node = 66 values
 - level 2 – at least $2*34$ nodes * 33 values/node = 2.244 values
 - level 3 – at least $2*34*34$ nodes * 33 values/node = 76.296 values
 - level 4 – at least $2*34*34*34$ nodes * 33 values/node = 2.594.064 values, which is greater than the number of existing values => this level does not appear in the tree

=> at most 4 levels in the tree

- after at most 4 block reads and a number of comparisons in main memory, it can be determined whether the value exists (the corresponding record's address can then be retrieved) or the search was unsuccessful

References

- [Ra02] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (3rd Edition), McGraw-Hill, 2002
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007,
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Ta13] ȚÂMBULEA, L., Curs Baze de date, Facultatea de Matematică și Informatică, UBB, 2013-2014
- [Si11] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts (6th Edition), McGraw-Hill, 2011
- [Kn76] KNUTH, D.E., Tratat de programare a calculatoarelor. Sortare și căutare. Ed. Tehnică, București, 1976
- [Ga09] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book (2nd Edition), Pearson Education, 2009